

CS2040S

Data Structures and Algorithms

AVL Trees

Puzzle of the Week:

100 prisoners. Every so often, one is chosen at random to enter a room with a light bulb. You can turn the light bulb on or off.

- **WIN** if one prisoner announces correctly that all have visited the room.
- **LOSE** if announcement is incorrect.

What if, initially, the state of the light is unknown, either on or off?

Housekeeping

PS4 Release 12 Feb 15:15

- Due: 18 Feb 23:59

Implement Scapegoat Trees!

- Beware: “It’s easy to introduce bugs”**
 - One of your TAs**

Take care to make sure you’re careful when implementing it.

Today's Plan

On the importance of being balanced

- Height-balanced binary search trees
- AVL trees
- Rotations
- Insertion Recap
- Deletion



Recap: Dictionary Interface

A collection of (key, value) pairs:

interface **IDictionary**

void insert(Key k, Value v)

insert (k,v) into table

Value search(Key k)

get value paired with k

Key successor(Key k)

find next key > k

Key predecessor(Key k)

find next key < k

void delete(Key k)

remove key k (and value)


boolean contains(Key k)

is there a value for k?

int size()

number of (k,v) pairs

Recap: Binary Search Trees



- Two children: $v.\text{left}$, $v.\text{right}$
- Key: $v.\text{key}$
- **BST Property:** all in left sub-tree $<$ key $<$ all in right sub-tree

Binary Search Tree

Modifying Operations: $O(h)$

- insert
- delete

Query Operations: $O(h)$

- search
- predecessor, successor
- findMax, findMin

Traversals: $O(n)$

The Importance of Being Balanced

Operations take $O(h)$ time

$$\log(n) - 1 \leq h \leq n$$

Key definition

A BST is balanced if $h = O(\log n)$

On a balanced BST: all operations run in $O(\log n)$ time.

The Importance of Being Balanced


How to get a balanced tree:

- Define a good property of a tree.
- Show that if the good property holds, then the tree is **balanced**.
- After every insert/delete, make sure the good property still holds. If not, fix it.



Invariant

AVL Trees [Adelson-Velskii & Landis 1962]




AVL Trees [Adelson-Velskii & Landis 1962]

Step 1: Define Invariant

- A node v is **height-balanced** if:

$$|v.\text{left.height} - v.\text{right.height}| \leq 1$$

Key definition



Height-Balanced Trees


Theorem:

A height-balanced tree with n nodes has **at most** height $h < 2\log(n)$.

- A height-balanced tree is balanced.

AVL Trees [Adelson-Velskii & Landis 1962]

Step 2: Show how to maintain height-balance



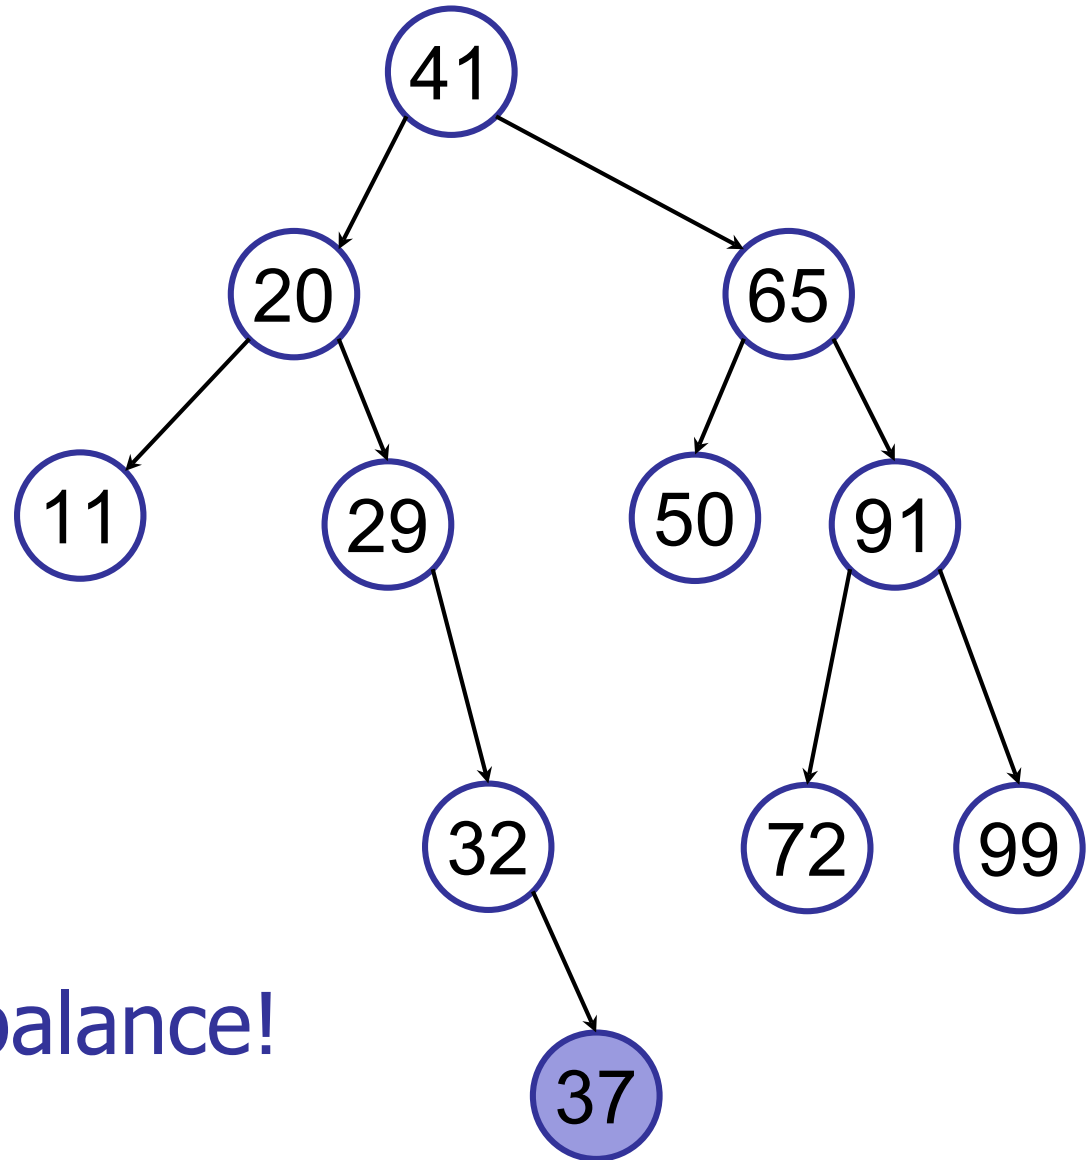
Inserting in an AVL Tree

Initially balanced

insert(37)

No longer balanced
after insertion!


Use rotations to rebalance!




Quick review: a rotation costs:

- ✓ 1. $O(1)$
- 2. $O(\log n)$
- 3. $O(n)$
- 4. $O(n^2)$
- 5. $O(2^n)$

Tree Rotations




Right
Rotation



$A < B < C < D < E$

Rotations maintain ordering of keys.
⇒ Maintains BST property.


Tree Rotations




Right
Rotation



Left
Rotation



Tree Rotations




After insert:

Use tree rotations to restore balance.

Height is out-of-balance by 1


Tree Rotations



A is **LEFT-heavy** if left sub-tree has larger height than right sub-tree.

A is **RIGHT-heavy** if right sub-tree has larger height than left sub-tree.


Tree Rotations



Use tree rotations to restore balance.

After insert, start at bottom, work your way up.


Tree Rotations



Assume **A** is the lowest node in the tree violating balance property.

Assume A is **LEFT-heavy**.

Tree Rotations (Left Heavy)




Assume **A** is the lowest node in the tree violating balance property.


Case 1: **B** is equi-height : $h(\text{L}) = h(\text{M})$

$$h(\text{R}) = h(\text{M}) - 1$$

Tree Rotations



Right
Rotation




right-rotate:

Case 1: **B** is equi-height : $h(\mathbf{L}) = h(\mathbf{M})$

$$h(\mathbf{R}) = h(\mathbf{M}) - 1$$

Tree Rotations (Left Heavy)




Assume **A** is the lowest node in the tree violating balance property.


Case 2: **B** is left-heavy : $h(\text{L}) = h(\text{M}) + 1$

$$h(\text{R}) = h(\text{M})$$

Tree Rotations



Right
Rotation




right-rotate:


Case 2: **B** is left-heavy: $h(\mathbf{L}) = h(\mathbf{M}) + 1$

$$h(\mathbf{R}) = h(\mathbf{M})$$

Tree Rotations



Right
Rotation




right-rotate:

Case 2: **B** is left-heavy: $h(\mathbf{L}) = h(\mathbf{M}) + 1$

$$h(\mathbf{R}) = h(\mathbf{M})$$

Tree Rotations (Left Heavy)




Assume **A** is the lowest node in the tree violating balance property.


Case 3: **B** is right-heavy : $h(\text{L}) = h(\text{M}) - 1$

$$h(\text{R}) = h(\text{L})$$

Tree Rotations



Right
Rotation




right-rotate:

Case 3: **B** is right-heavy: $h(\mathbf{L}) = h(\mathbf{M}) - 1$

$$h(\mathbf{R}) = h(\mathbf{L})$$

Tree Rotations



Let's do something
first before we
`right-rotate(A)`


(Reduce it to a
problem we have
already solved!)

`right-rotate:`

Case 3: **B** is right-heavy: $h(\mathbf{L}) = h(\mathbf{M}) - 1$

$$h(\mathbf{R}) = h(\mathbf{L})$$


Tree Rotations




Left-rotate B

After left-rotate B: **A** and **C** still out of balance.

Tree Rotations



Tree Rotations



After right-rotate A: all in balance.

Rotations

Summary:

If v is out of balance and left heavy:

1. $v.\text{left}$ is balanced: $\text{right-rotate}(v)$
2. $v.\text{left}$ is left-heavy: $\text{right-rotate}(v)$
3. $v.\text{left}$ is right-heavy: $\text{left-rotate}(v.\text{left})$
 $\text{right-rotate}(v)$

If v is out of balance and right heavy:

Symmetric three cases....

How many rotations do you need after an insertion (in the worst case)?


1. 1
2. 2
3. 4
4. $\log(n)$
5. $2\log(n)$
6. n

How many rotations do you need after an insertion (in the worst case)?

- 1. 1
- ✓ 2. 2
- 3. 4
- 4. $\log(n)$
- 5. $2\log(n)$
- 6. n

Question:
Why isn't it $2\log(n)$?


How many rotations?



Case 2: **B** is left-heavy

Insert increased heights by 1.

How many rotations?




Case 2: **B** is left-heavy

Rotation reduces root height by 1.

(Everything higher in tree is unchanged!)


How many rotations?



Case 3: **B** is right-heavy

Rotation reduces root height by 1.

How many rotations?




Case 1: **B** is balanced


Rotation does *not* reduce height by 1.

Challenge: figure out why this is okay!

Root height does not decrease!



Right
Rotation




right-rotate:

Case 1: **B** is equi-height : $h(\mathbf{L}) = h(\mathbf{M})$

$$h(\mathbf{R}) = h(\mathbf{M}) - 1$$


How did we get here?



If the tree was balanced before the insert...

... no possible insert could have increased the height of both **L** and **M**.

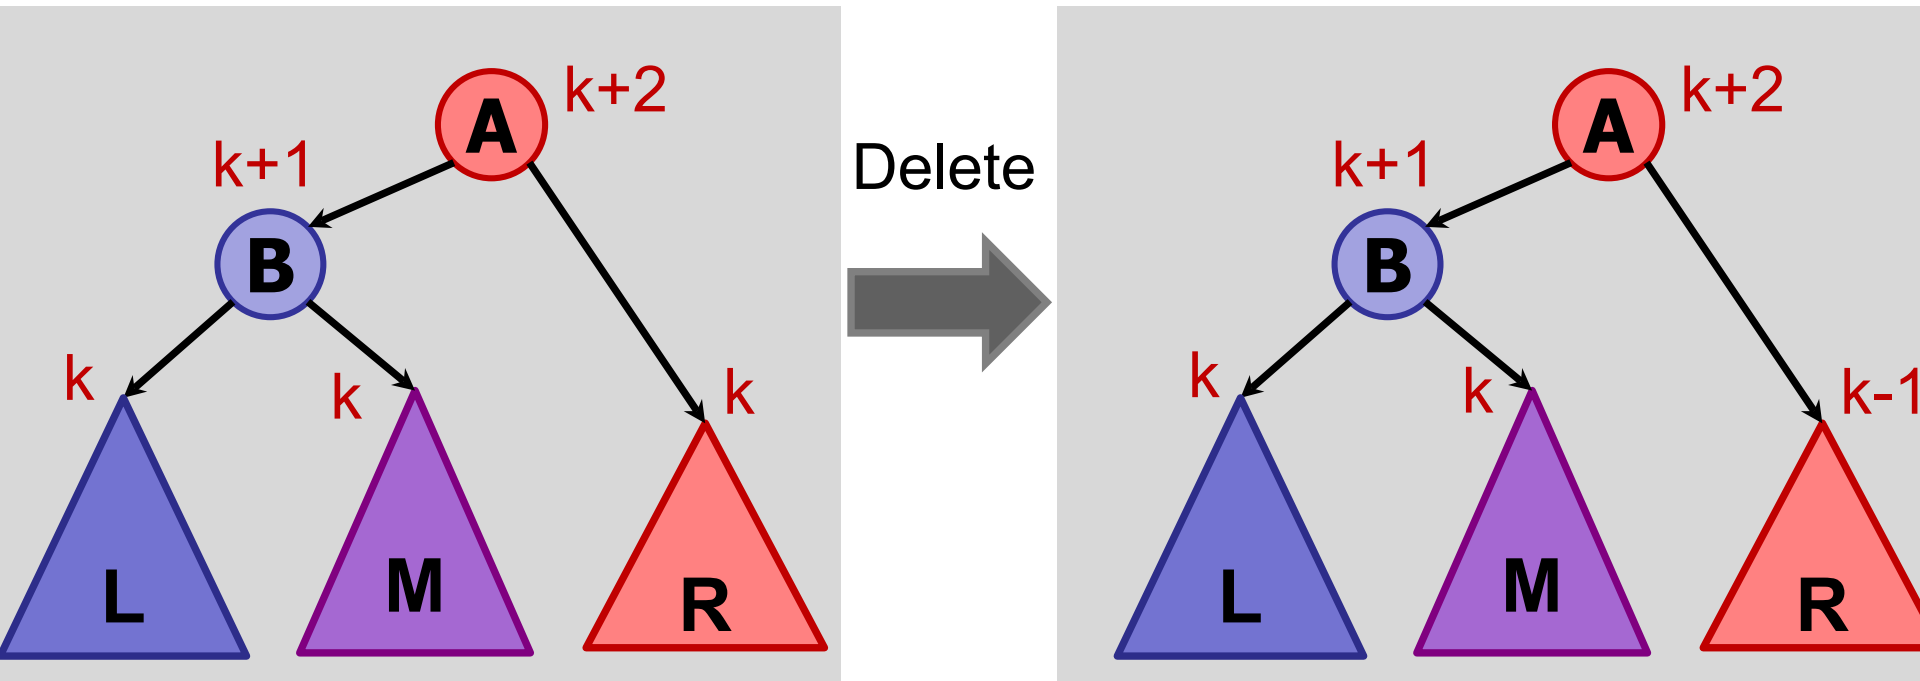
Root height does not decrease!



Why did we cover Case 1?


We need this case for deletes...

How did we get here?



Delete in tree R unbalances the tree...

How did we get here?



Delete in tree R unbalances the tree...

And after the rotation to fix it...

Insert in AVL Tree

Summary:


- Insert key in BST.
- Walk up tree:
 - At every step, check for balance.
 - If out-of-balance, use rotations to rebalance and return.

Key observation:

- Only need to fix *lowest* out-of-balance node.
- Only need at most two rotations to fix.



Example

insert(23)





Example

insert(23)




Example

right-rotate(29)





Example

insert(55)





Example

insert(55)





Example

left-rotate(50)



Example

right-rotate(65)




Binary Search Tree

delete(v)

Three cases:

1. No children
2. 1 child
3. 2 children



Binary Search Tree


delete(v)

1. If **v** has two children, swap it with its successor.
2. Delete node v from binary tree (and reconnect children).
3. For every ancestor of the deleted node:
 - Check if it is height-balanced.
 - If not, perform a rotation.
 - Continue to the root.

Deletion may take up to $O(\log(n))$ rotations.

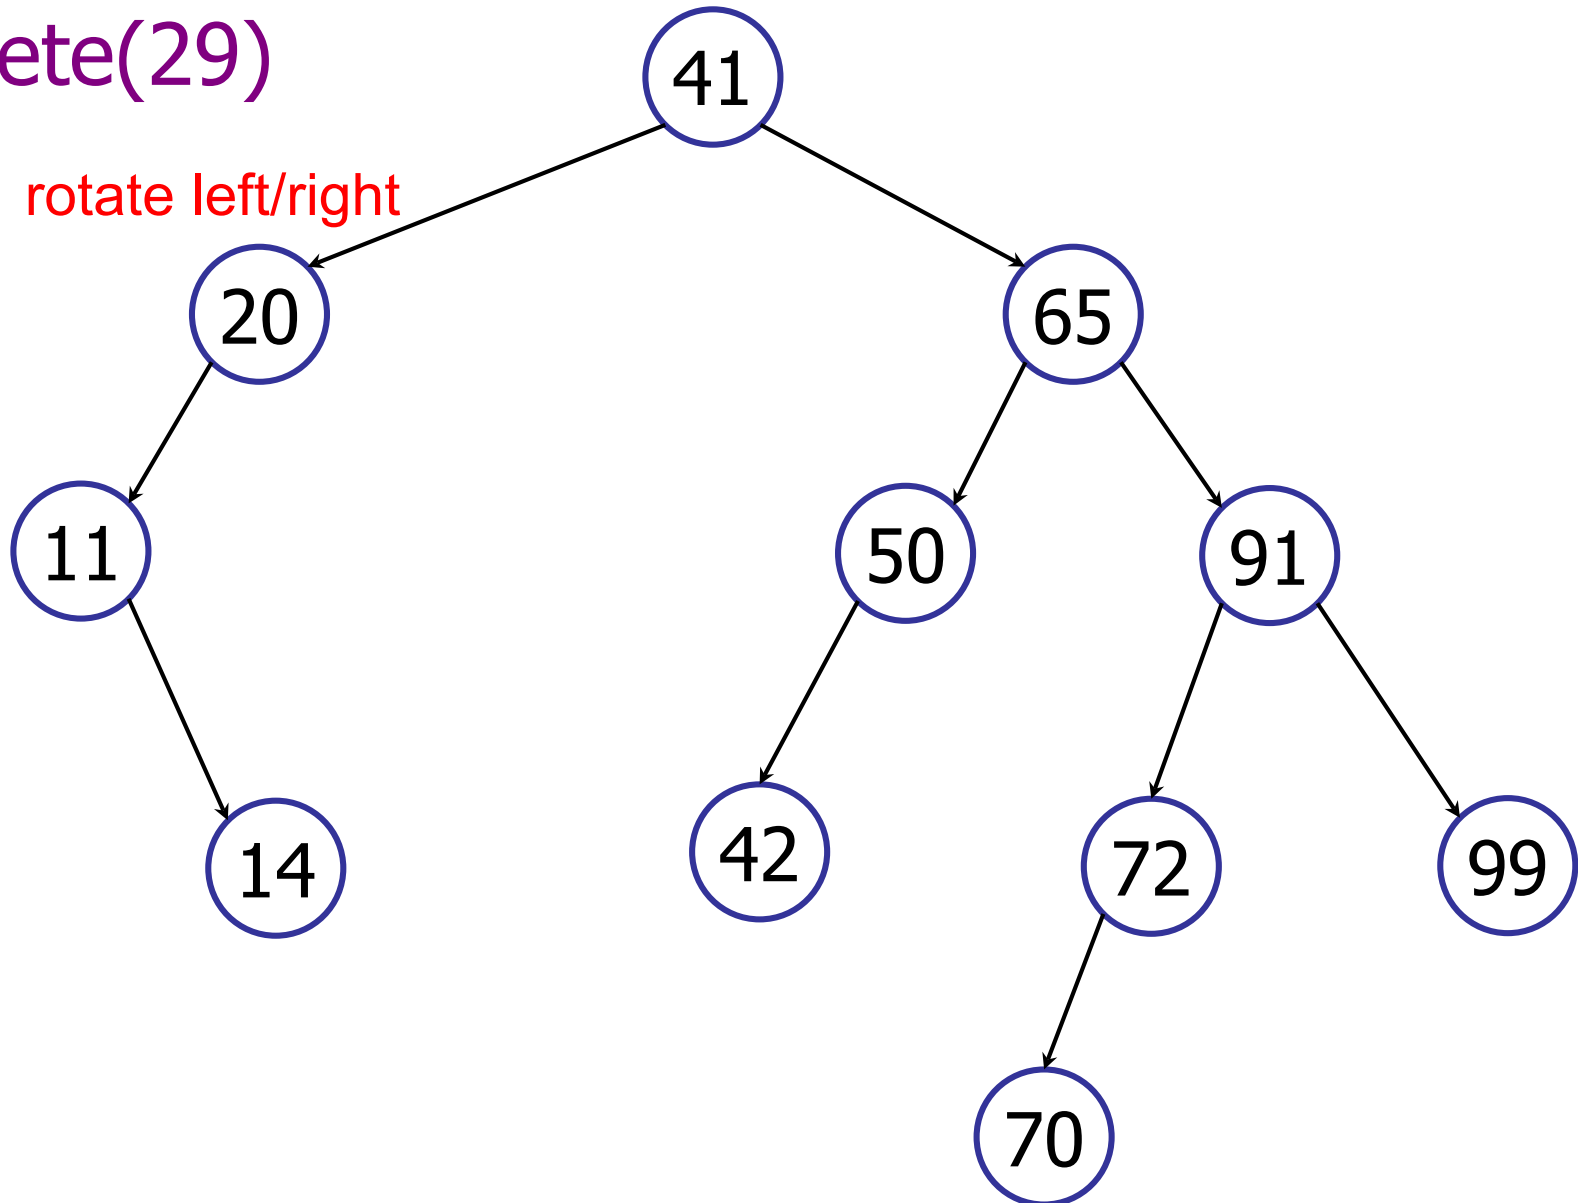
Binary Search Tree

delete(29)



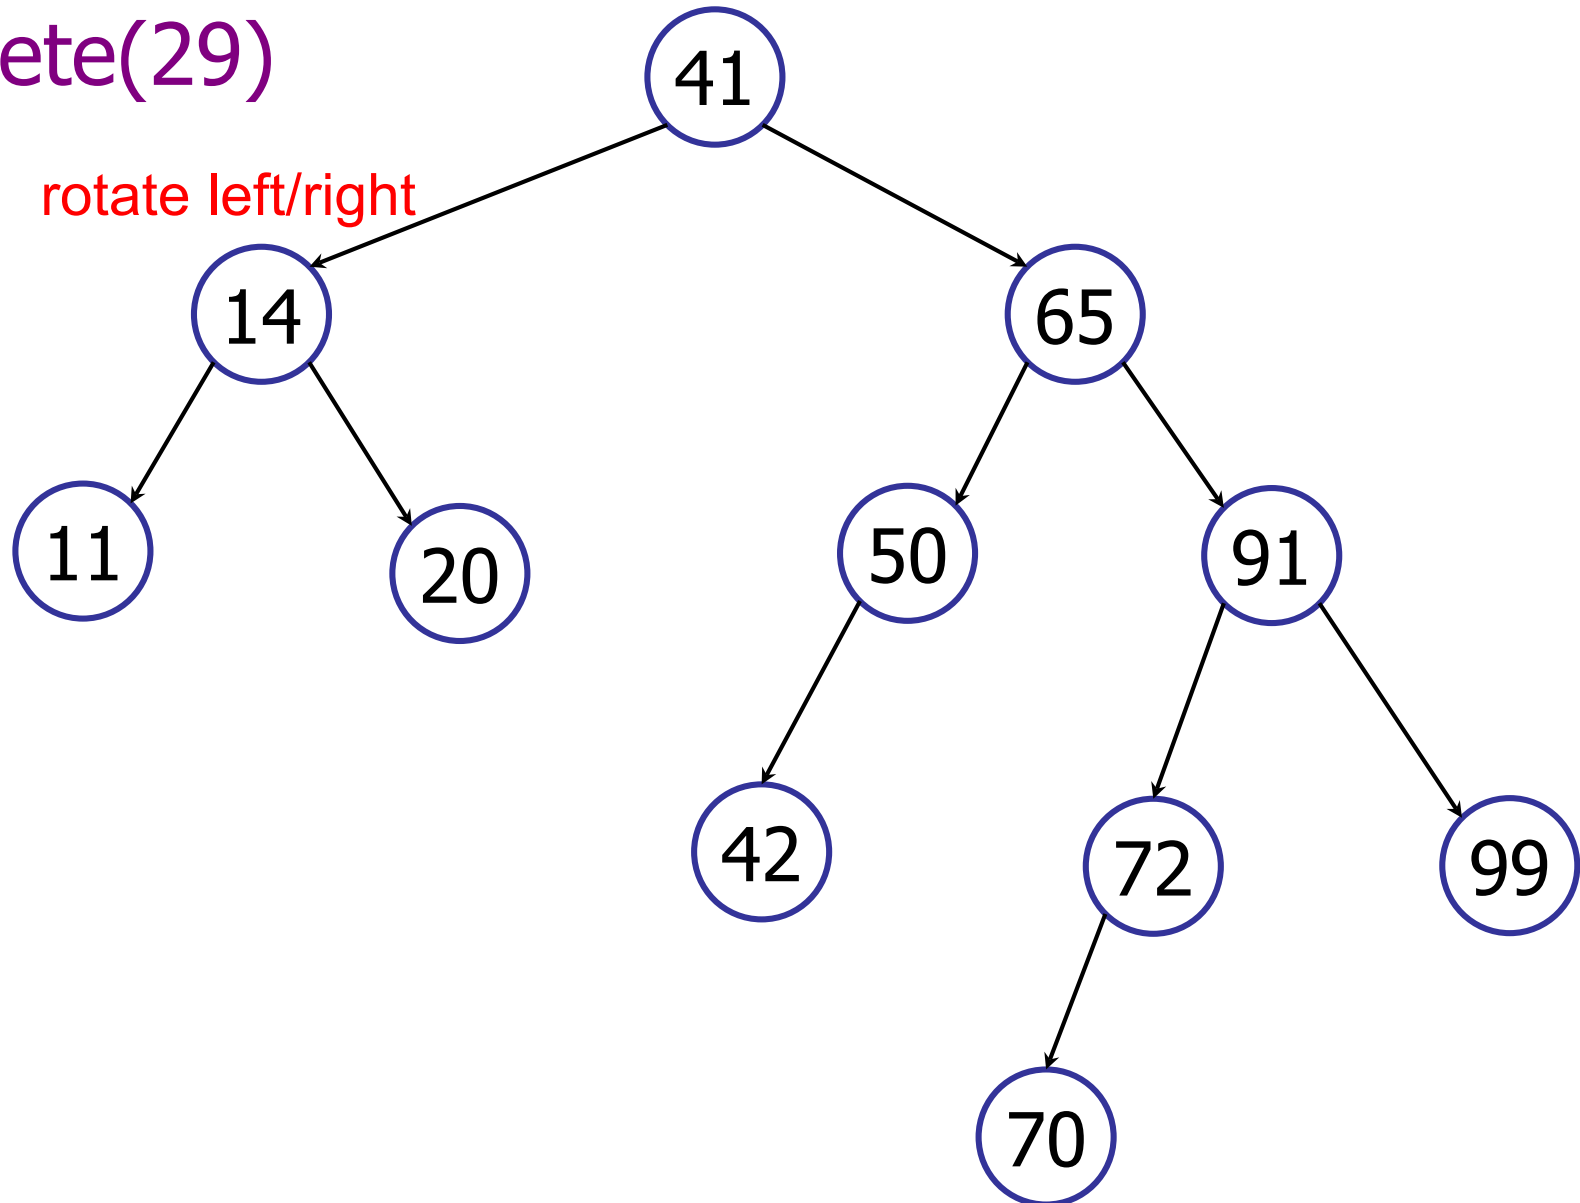
Binary Search Tree

delete(29)




Binary Search Tree

delete(29)




Binary Search Tree

delete(29)



Binary Search Tree

delete(29)



How many rebalances?

Why are two rotations not enough?

- Delete reduced height.
- Rotations (to rebalance) reduce height!

Key observation:

- Rebalancing does not “undo” the change in height caused by deletion.

Delete in AVL Tree

Summary:

- Delete key from BST.
- Walk up tree:
 - At every step, check for balance.
 - If out-of-balance, use rotations to rebalance.
 - Continue to root.

Key observation:

- It is *not* sufficient to only fix lowest out-of-balance node in tree.

Every insertion requires 1 or 2 rotations?

1. Yes
- ✓ 2. No
3. I don't know

A tree is **balanced** if every node's children differ in height be at most 1?


- ✓ 1. Yes
- 2. No
- 3. I don't know

A tree is **balanced** if every node either has two children or zero children?

1. Yes
- ✓ 2. No
3. I don't know


A tree is balanced if every node either has two children or zero children?

1. Yes
- ✓ 2. No
3. I don't know



A tree is height-balanced if:

For every node, the number of keys in its heavier sub-tree is at most twice the number of keys in its lighter sub-tree.

1. Yes
-  2. No, but it is balanced.
3. No.
4. I don't know

Using rotations, you can create every possible “tree shape.”

- ✓ 1. True
- 2. False
- 3. I don't know

AVL Trees: Other potential modifications

What if you do not remove deleted nodes?

- Mark a node “deleted” and leave it in the tree.

Logical deletes:

- Performance degrades over time.
- Clean up later? (Amortized performance...)

AVL Trees: Other potential modifications

What if you do not want to store the height in every node?

- Only store difference in height from parent.

Next Week

Even more trees! (Monday)

- Other augmentations
- Examples of other forms of trees

Hashing! (Wednesday)

- Introduction to hashing