

CS2040S

Data Structures and Algorithms

Bellman-Ford


Last Time

Single Source Shortest Paths (SSSP):

- Dijkstra
 - SSSP on non-negatively weighted graphs


Dijkstra's Algorithm

Edges with negative weights?



Dijkstra's Algorithm


Edges with negative weights?



Step 1: Remove A.
Relax A.
Mark A done.

Dijkstra's Algorithm

Edges with negative weights?



Step 1: Remove A.
Relax A.
Mark A done.

...


Step 4: Remove B.
Relax B.
Mark B done.

Oops: We need to
update A.

Dijkstra's Algorithm

Can we reweight?

e.g.: $\text{weight} += 10$




Can we reweight the graph?

1. Yes.
2. Only if there are no negative weight cycles.
- ✓ 3. No.

Dijkstra's Algorithm

Can we reweight?




Path S-B-A: 1

Path S-A: 4

Dijkstra's Algorithm

Can we reweight?




Path S-B-A: 21

Path S-A: 14


Negative Cycles:

What if edges have negative weight?



Negative Cycles:

What if edges have negative weight?



$d(S,C)$ is infinitely negative!

Today

Single Source Shortest Paths (SSSP):

- Bellman Ford
 - SSSP on negative edge graphs
 - Negative Cycle Detection


Some graph techniques

Recall: Path Relaxation

Let's quickly refresh path relaxation, we're going to be doing a lot of that today.


Recall: Path Relaxation

```
relax(int u, int v){  
    if (dist[v] > dist[u] + weight(u,v))  
        dist[v] = dist[u] + weight(u,v);  
}
```




Recall: Path Relaxation

Let's try running path relaxation again.




Recall: Path Relaxation

Our distance estimates.




Recall: Path Relaxation

Let's say we ran `relax()` based on the edges we have, in some arbitrary ordering.



Recall: Path Relaxation


$\text{relax}(S, A)$



Recall: Path Relaxation


$\text{relax}(S, A)$

Reduced from ∞ to 3.



Recall: Path Relaxation


$\text{relax}(A, C)$



Recall: Path Relaxation


$\text{relax}(A, C)$

Reduced from ∞ to $3 + 4 = 7$.



Recall: Path Relaxation


$\text{relax}(A, B)$



Recall: Path Relaxation


$\text{relax}(A, B)$

Reduced from ∞ to $3 + 1 = 4$.



Recall: Path Relaxation


$\text{relax}(S, B)$



Recall: Path Relaxation


$\text{relax}(S, B)$

No change!



Recall: Path Relaxation


$\text{relax}(B, C)$



Recall: Path Relaxation

$\text{relax}(B, C)$


Reduced from 7 to $4 + 2 = 6$.



Shortest Paths

for (edge **e** : graph)

 relax(**e**)




Shortest Paths

for (edge **e** : graph)


 relax(**e**)

Let's say the order in which we iterate over the edges is not determined by us.



Does this algorithm work?
for every edge e : relax(e)

1. Yes
2. Sometimes
3. Never




Shortest Paths

for (edge **e** : graph)

 relax(**e**)

What if the ordering
was:

(A, C)
(A, B)
(B, C)
(S, A)
(S, B)




What happens if we ran this for a single round? What is the distance estimate of A?

- ✓ 1. 3
2. ∞

What if the ordering was:

- (A, C)
- (A, B)
- (B, C)
- (S, A)
- (S, B)




What happens if we ran this for a single round? What is the distance estimate of B?

1. 3
2. 4
- ☒ 3. 5
4. ∞

What if the ordering was:

- (A, C)
- (A, B)
- (B, C)
- (S, A)
- (S, B)



What happens if we ran this for a single round? What is the distance estimate of C?

1. 7

2. 6

3. ∞

What if the ordering was:


(A, C)

(A, B)

(B, C)

(S, A)

(S, B)



Shortest Paths


for (edge **e** : graph)

 relax(**e**)

After 1 round:

What if the ordering
was:

- (A, C)
- (A, B)
- (B, C)
- (S, A)
- (S, B)




Can we say that the nodes that are one hop away from source S have correct distance estimates?

1. Yes
2. No
3. Narp

What if the ordering was:

- (A, C)
- (A, B)
- (B, C)
- (S, A)
- (S, B)



Shortest Paths

for (edge **e** : graph)


 relax(**e**)

After 1 round:

Node B's distance estimate is still not correct!

What if the ordering was:

- (A, C)
- (A, B)
- (B, C)
- (S, A)
- (S, B)



Shortest Paths


for (edge **e** : graph)

 relax(**e**)

So this alone clearly doesn't work. But can we at least say we're making some progress?


Idea:

Let's consider some general directed, weighted graph.




Idea:

Consider the shortest path tree of the graph:



Idea:

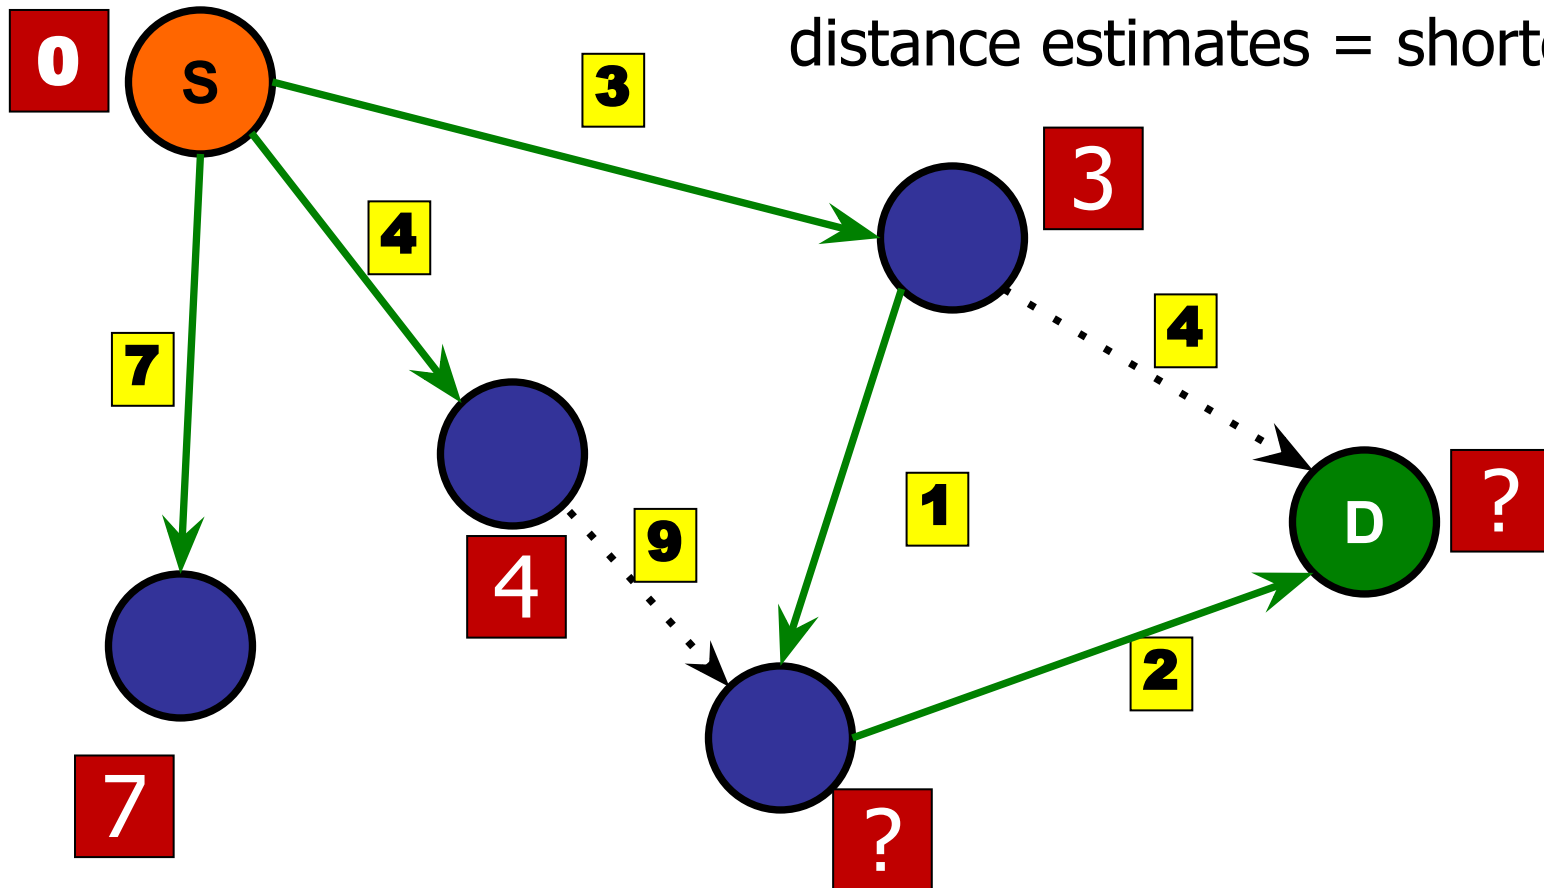
What can we say about the distance estimates after one round of path relaxation over the edges?



Idea:

After 1 round of relaxation, the nodes that are 1 hop away on the shortest path tree, have their

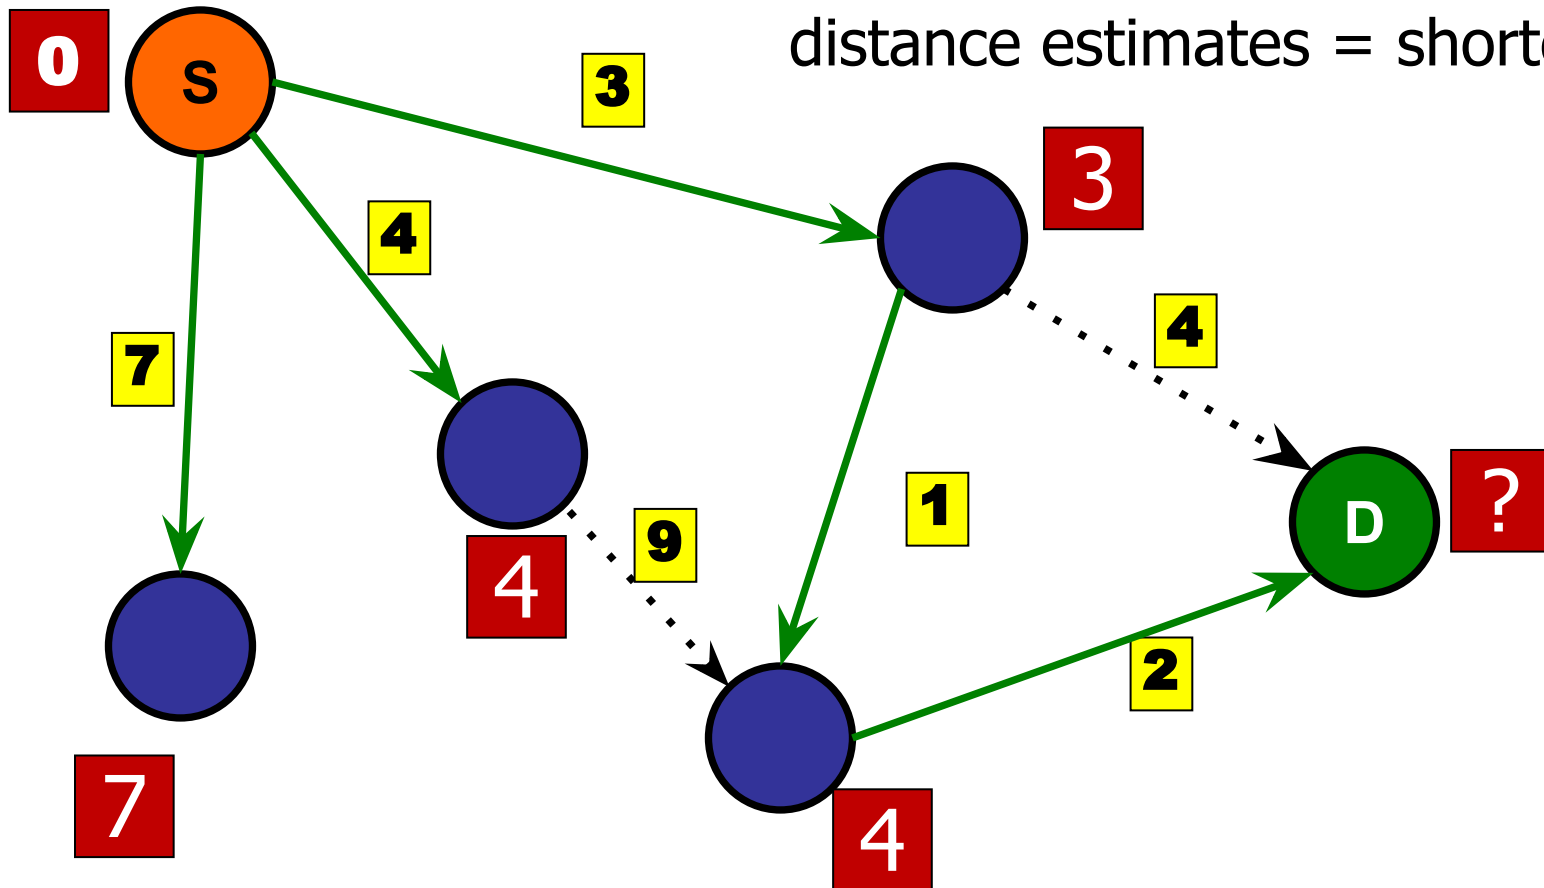
distance estimates = shortest dist



Idea:

After 2 rounds of relaxation, the nodes that are 2 hop away on the shortest path tree, have their

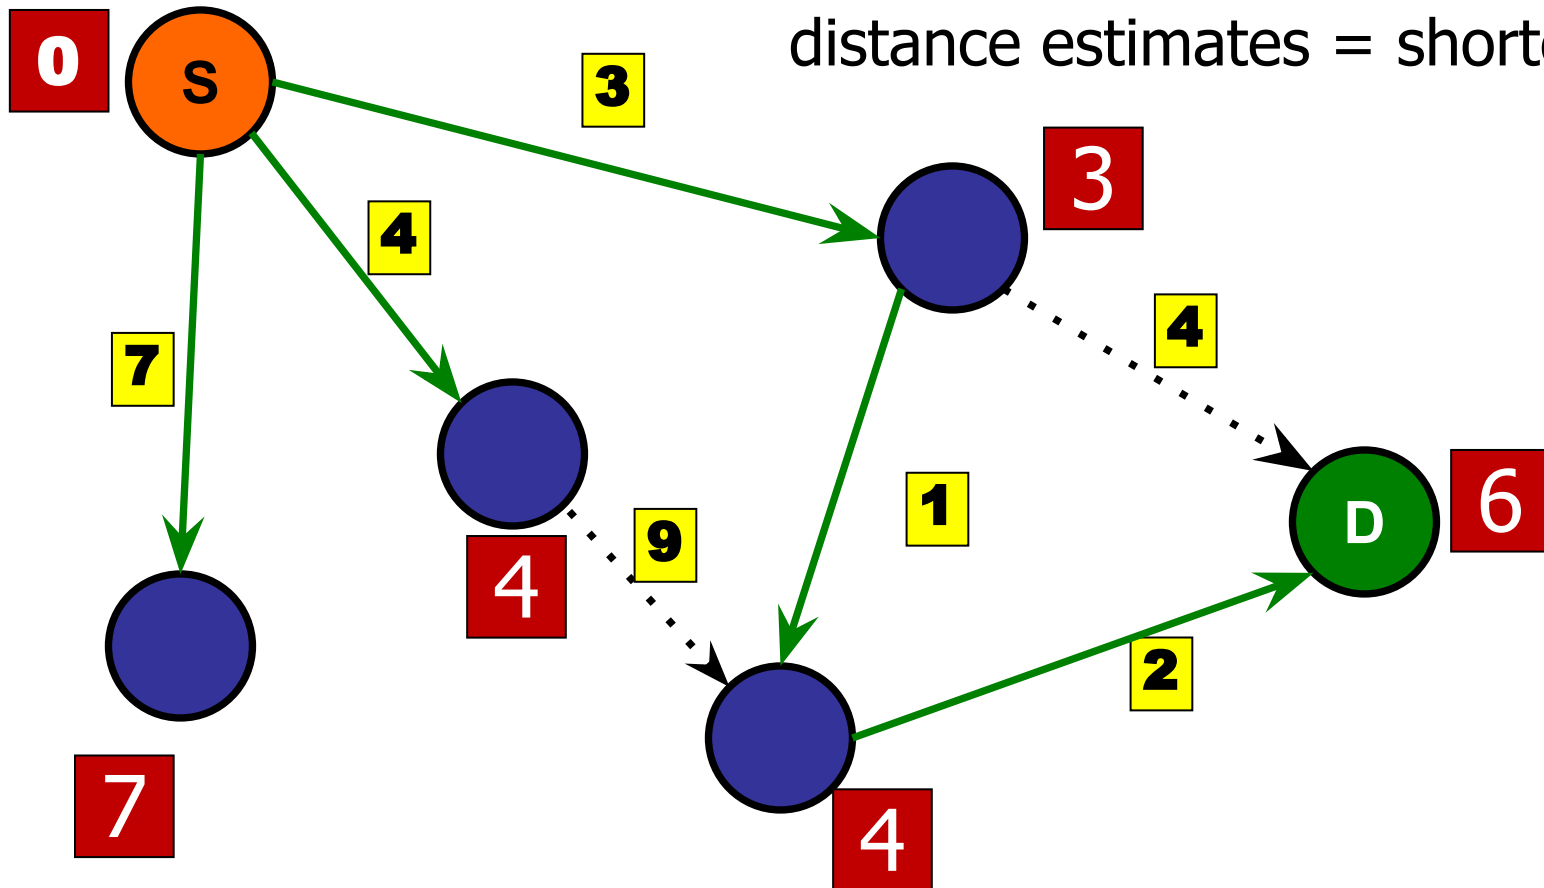
distance estimates = shortest dist



Idea:


After 3 rounds of relaxation, the nodes that are 3 hop away on the shortest path tree, have their

distance estimates = shortest dist




Idea:

To be clear: It takes **at most** i rounds to compute the correct distance that are i hops away on the shortest path tree



Idea:

Corollary: Since every node has to be at most $|V| - 1$ hops away from the source node s , we just need to run $|V| - 1$ rounds.



Bellman-Ford:

Pseudocode:

Set up distance estimate array `dist`

for $|V| - 1$ iterations:

 for edge (u, v) in the graph G :

`relax(dist, u, v)`

What is the time complexity of the given algorithm?

Pseudocode:

Set up distance estimate array `dist`

for $|V| - 1$ iterations:

for edge (u, v) in the graph G :

`relaxed = relax(dist, u, v)`

1. $O(V + E)$

✓ 2. $O(VE)$

3. $O(V^2)$

4. $O(E^2)$

Bellman-Ford

Claim:

If after a round, the distance estimates don't change, we have found the shortest distances for all nodes.

Bellman-Ford

If after a round, the distance estimates don't change, we have found the shortest distances for all nodes.

Intuition:

Let's say **before** we ran a round of relaxations, and we started with distance array D_1 .

It didn't change **after** the round of relaxations.

So even if we ran even more iterations (up until all $|V| - 1$ of them), nothing will change.

Bellman-Ford

If after a round, the distance estimates don't change, we have found the shortest distances for all nodes.

Intuition:

Let's say **before** we ran a round of relaxations, and we started with distance array D_1 .

It didn't change **after** the round of relaxations.

So even if we ran even more iterations (up until all $|V| - 1$ of them), nothing will change.

Early termination!

Bellman-Ford:

Pseudocode:

Set up distance estimate array `dist`

for $|V| - 1$ iterations:

 for edge (u, v) in the graph G :


`relaxed` \leftarrow `relax(dist, u, v)`

 if not `relaxed`: // no estimates have changed

 break

Bellman-Ford

What if edges have negative weight?



Bellman-Ford

What if edges have negative weight?

Assume ordering was:


(A, C)

(C, B)

(A, B)

(S, A)

(S, B)



Bellman-Ford

What if edges have negative weight?

Assume ordering was:


(A, C)

(C, B)

(A, B)

(S, A)

(S, B)



After 1 round.

Bellman-Ford

What if edges have negative weight?

Assume ordering was:


(A, C)

(C, B)

(A, B)

(S, A)

(S, B)



After 2 rounds.

Bellman-Ford

What if edges have negative weight?

Assume ordering was:


(A, C)

(C, B)

(A, B)

(S, A)

(S, B)



After 3 rounds. No changes already.

Bellman-Ford

What if edges have negative weight?

Assume ordering was:


(A, C)

(C, B)

(A, B)

(S, A)

(S, B)



After 3 rounds. Shortest distances found!

Bellman-Ford

What if edges have negative weight?

Assume ordering was:


(A, C)

(C, B)

(A, B)

(S, A)

(S, B)



No problem!

Bellman-Ford

What if the graph has a negative cycle?

Assume ordering was:


(A, C)

(C, B)

(B, A)

(S, A)

(S, B)



Bellman-Ford

What if the graph has a negative cycle?

Assume ordering was:


(A, C)

(C, B)

(B, A)

(S, A)

(S, B)



After 1 rounds.

Bellman-Ford

What if the graph has a negative cycle?

Assume ordering was:


(A, C)

(C, B)

(B, A)

(S, A)

(S, B)



After 2 rounds.

Bellman-Ford

What if the graph has a negative cycle?

Assume ordering was:


(A, C)

(C, B)

(B, A)

(S, A)

(S, B)



After 3 rounds.

Bellman-Ford

What if the graph has a negative cycle?

Assume ordering was:


(A, C)

(C, B)

(B, A)

(S, A)

(S, B)



After ???? rounds.

Bellman-Ford

What if the graph has a negative cycle?

Assume ordering was:


(A, C)

(C, B)

(B, A)

(S, A)

(S, B)



$d(S, C)$ is infinitely negative!

Bellman-Ford

What if the graph has a negative cycle?

Assume ordering was:


(A, C)

(C, B)

(B, A)

(S, A)


(S, B)



Notice here that we don't even have a cycle made of all negative edges.


Detecting negative cycles

We know that any shortest paths should be found after $|V| - 1$ iterations.




If there is a negative cycle, what happens to the distance estimates on the $|V|^{th}$ iteration?

1. Estimates remain unchanged
2. ✓ Some estimates will go down
3. Some estimates will go up



Detecting negative cycles

We know that any shortest paths should be found after $|V| - 1$ iterations.




Run one more iteration, if any distance estimates go down, we know there is a negative cycle

Shortest Paths (Recall)

Key idea: triangle inequality

$$\delta(S, C) \leq \delta(S, A) + \delta(A, C)$$


(Side Quiz: Does this also hold if our edge weights are negative?)



So Far:

- **Unweighted** graph
 - BFS
 - $O(V + E)$
- **Weighted** graphs with non-negative edges.
 - Dijkstra
 - $O(E \log V)$
- **Weighted** graphs with no negative cycles.
 - Bellman-Ford
 - $O(VE)$


Not the end of the story:



Previously:

Let's say we ran `relax()` based on the edges we have, in some arbitrary ordering.

What if we had control over this?



Changing the Ordering:

```
for (edge e : graph)
```

```
    relax(e)
```




```
for (edge e : graph)
```

```
    relax(e)
```



Bellman-Ford on DAG


What happens if the graph is a DAG?



Bellman-Ford on DAG

What happens if the graph is a DAG?


Toposort it first!



Bellman-Ford on DAG

What happens if the graph is a DAG?


Toposort it first!



Toposorted order: S, A, C, B

Bellman-Ford on DAG

In topo-sort order:



Toposorted order: S, A, C, B

Bellman-Ford On DAG:

Pseudocode:

Set up **distance estimate** array **dist**

Get **toposorted list** of nodes **topo_list**

for **u** in **topo_list**: (from first to last)

 for neighbour **v** in **u.neighbour_list**:

relax(**dist**, **u**, **v**)

What is the time complexity of this algorithm?

- ✓ 1. $O(V + E)$
 - 2. $O(V^2)$
 - 3. $O(VE)$
 - 4. $O(E^2)$
- Set up **distance estimate** array **dist**
Get **toposorted** list of nodes **topo_list**
for **u** in **topo_list**: (from first to last)
 for neighbour **v** in **u.neighbour_list**:
 relax(**dist**, **u**, **v**)

Bellman-Ford On DAG:

Pseudocode:

$O(V + E)$

Set up distance estimate array `dist`

Get toposorted list of nodes `topo_list`

for `u` in `topo_list`: (from first to last)


 for neighbour `v` in `u.neighbour_list`:

`relax(dist, u, v)`

$O(V + E)$


Bellman-Ford on DAG

In topo-sort order: S, A, C, B




Bellman-Ford on DAG

In topo-sort order: S, A, C, B




Bellman-Ford on DAG

In topo-sort order: S, A, C, B




Bellman-Ford on DAG

In topo-sort order: S, A, C, B



Bellman-Ford on DAG

In topo-sort order: S, A, C, B



So Far:

- **Unweighted** graph
 - BFS
 - $O(V + E)$
 -
- **Weighted** graphs with non-negative edges.
 - Dijkstra
 - $O(E \log V)$
- **Weighted** graphs with no negative cycles.
 - Bellman-Ford
 - $O(VE)$ on general
 - $O(V + E)$ with toposort on DAG

Bonus: Active Research

What about if we could randomise?

Jeremy Fineman in STOC' 2024 showed an algorithm that runs in $O(EV^{8/9})$ time with high probability.

[Conference](#) [Proceedings](#) [Upcoming Events](#) [Authors](#) [Affiliations](#) [Award Winners](#)

Home > Conferences > STOC > Proceedings > STOC 2024 > Single-Source Shortest Paths with Negative Real Weights in $\tilde{O}(mn^{8/9})$ Time

RESEARCH-ARTICLE ✕ in

Single-Source Shortest Paths with Negative Real Weights in $\tilde{O}(mn^{8/9})$ Time

Author: [Jeremy T. Fineman](#) [Authors Info & Claims](#)

STOC 2024: Proceedings of the 56th Annual ACM Symposium on Theory of Computing • Pages 3 - 14
<https://doi.org/10.1145/3618260.3649614>

Published: 11 June 2024 [Publication History](#)

0 636

STOC 2024: Proceedings

Abstract

Today

Single Source Shortest Paths (SSSP):

- Bellman Ford
 - SSSP on negative edge graphs
 - Negative Cycle Detection

Some graph techniques

Technique: Graph Modifications

Now that we've talked about single-source shortest paths. Let's think about:

Technique: Graph Modifications

Now that we've talked about single-source shortest paths. Let's think about:

What if we had multiple sources, and we just want the shortest path to/from one of these sources?

Technique: Graph Modifications


Now that we've talked about single-source shortest paths. Let's think about:

What if we had multiple sources, and we just want the shortest path to/from **any** of these sources?

E.g. we have many fire stations. We just want the closest one to reach the target as soon as possible.

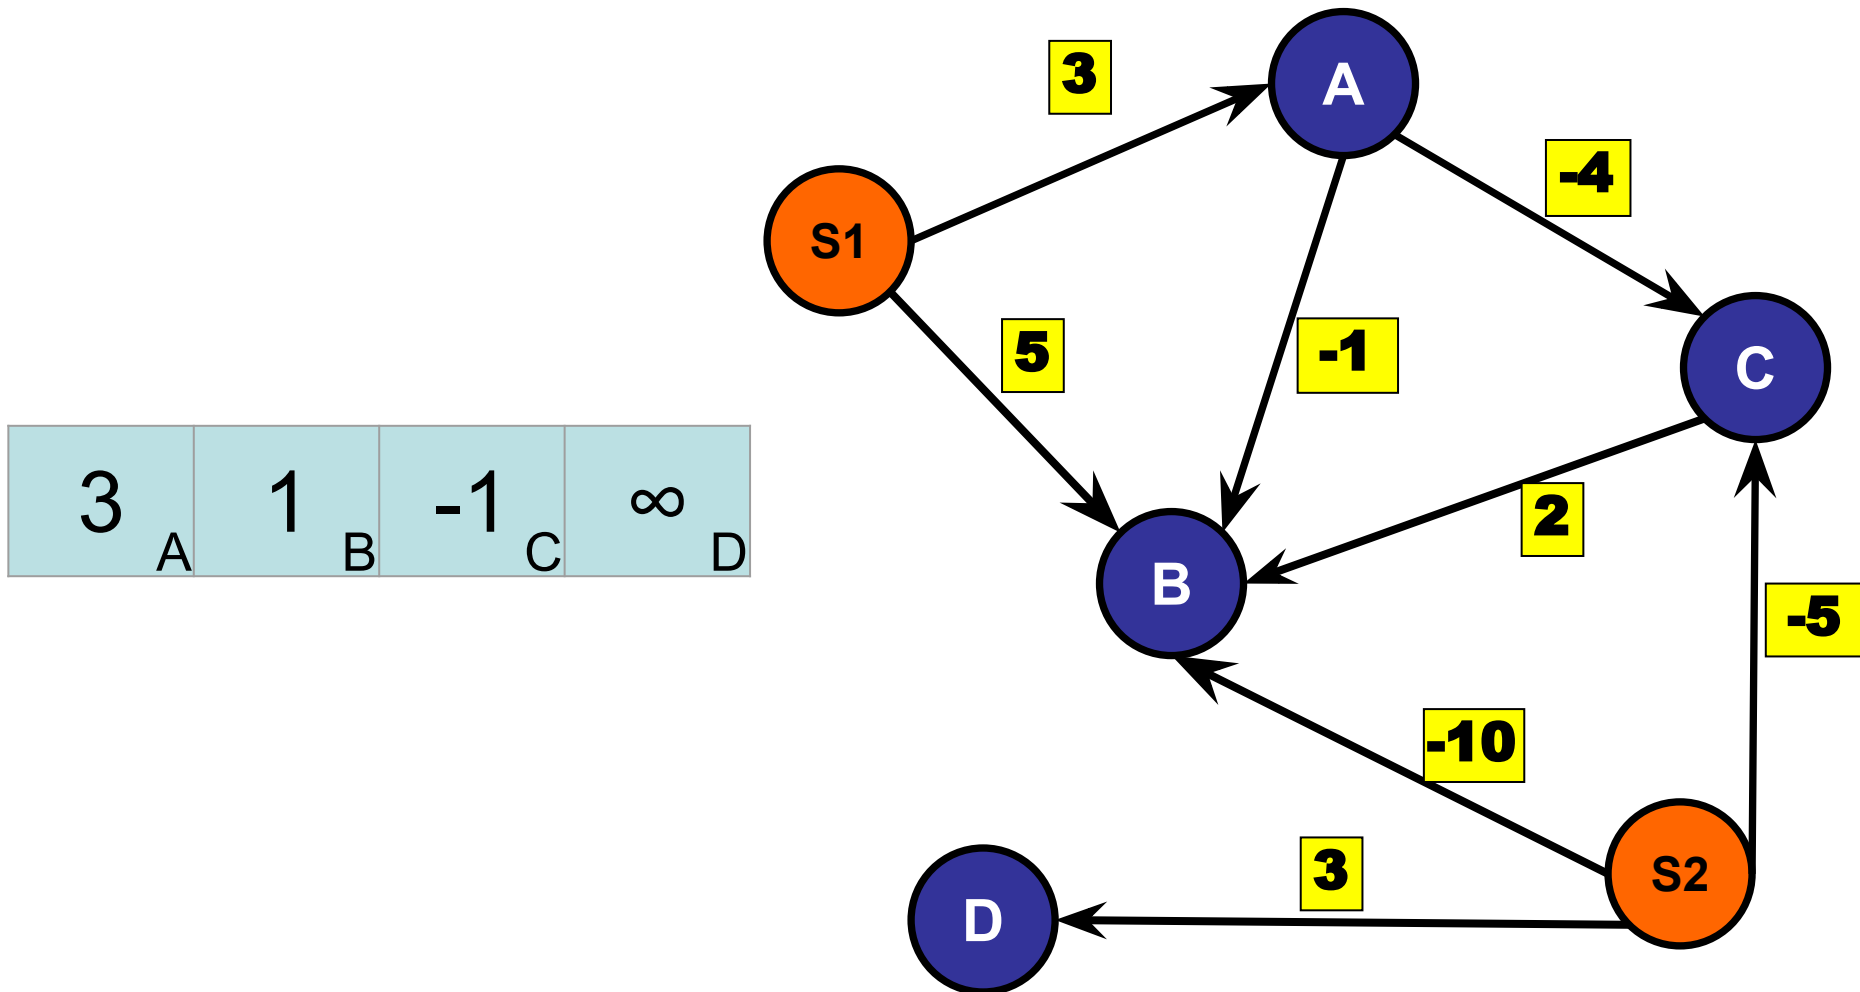
Technique: Graph Modifications

E.g. we have two sources s1 and s2.



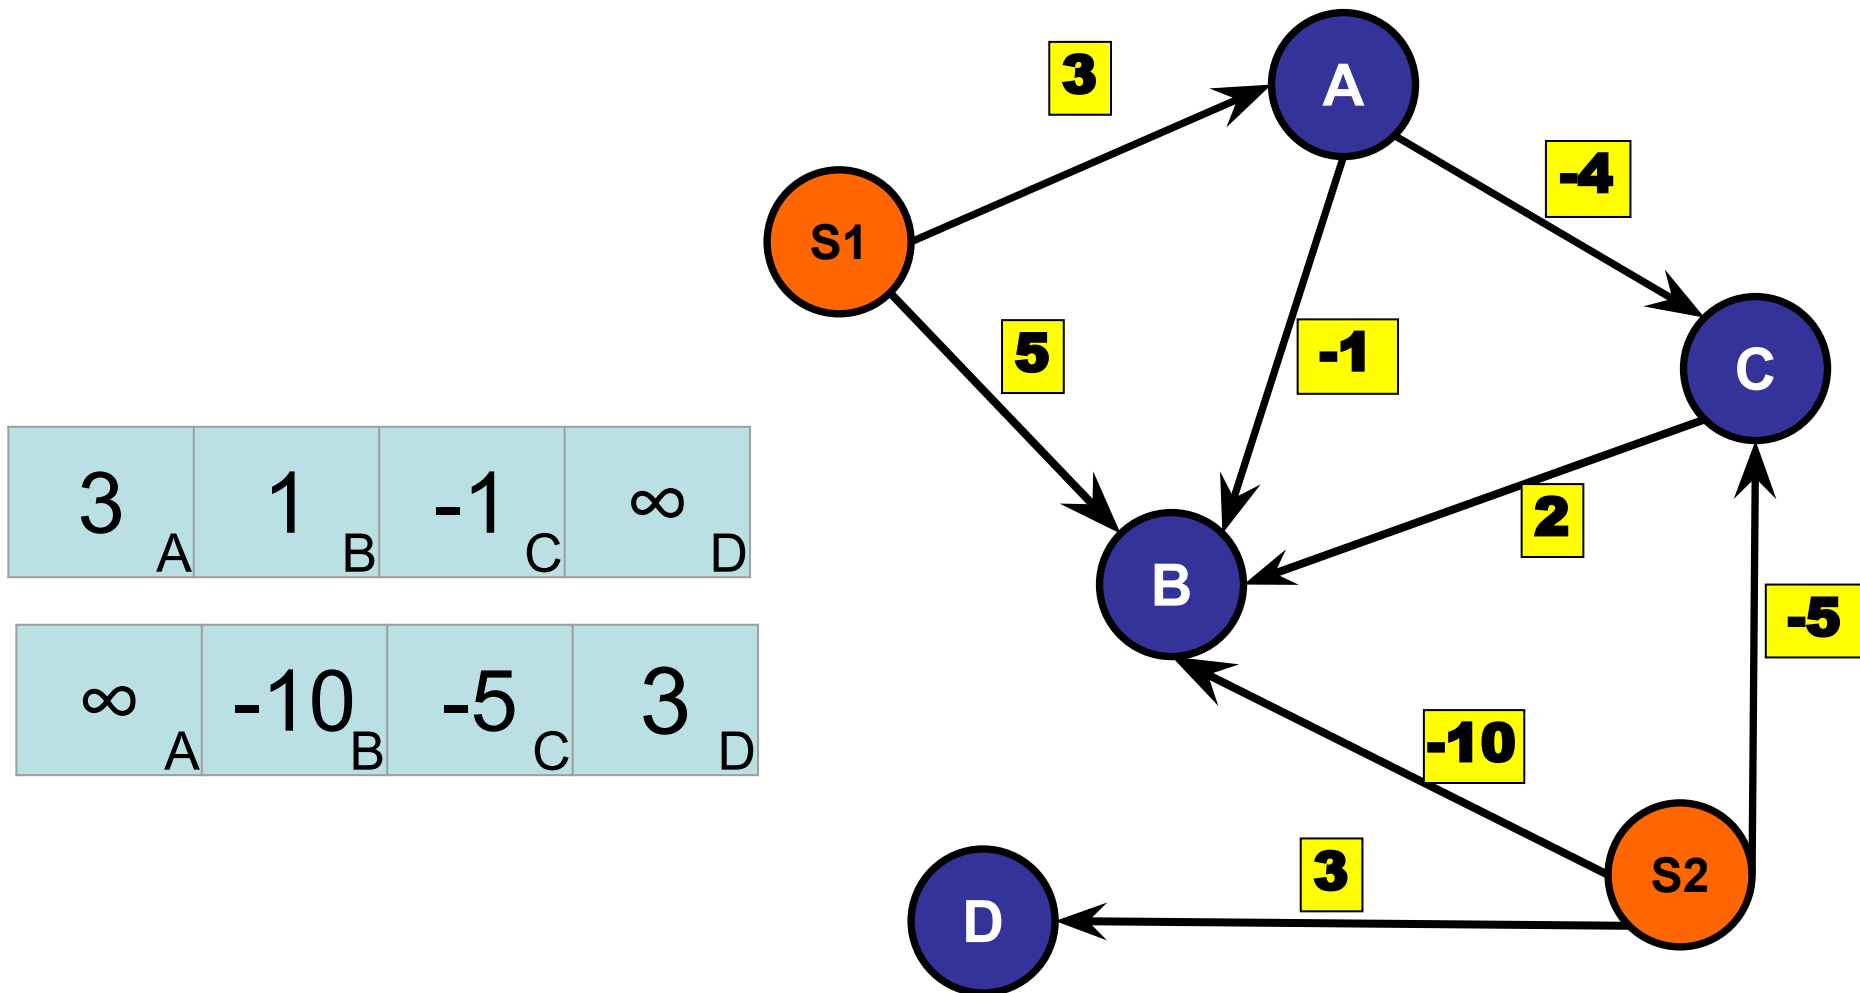
Technique: Graph Modifications

Obvious solution: Run SSSP once from S1.



Technique: Graph Modifications

Obvious solution: Run SSSP once from S1.
Run SSSP again from S2.




Technique: Graph Modifications

Obvious solution: Run SSSP once from S1.
Run SSSP again from S2.

Output minimum of both.

3	1	-1	∞
A	B	C	D

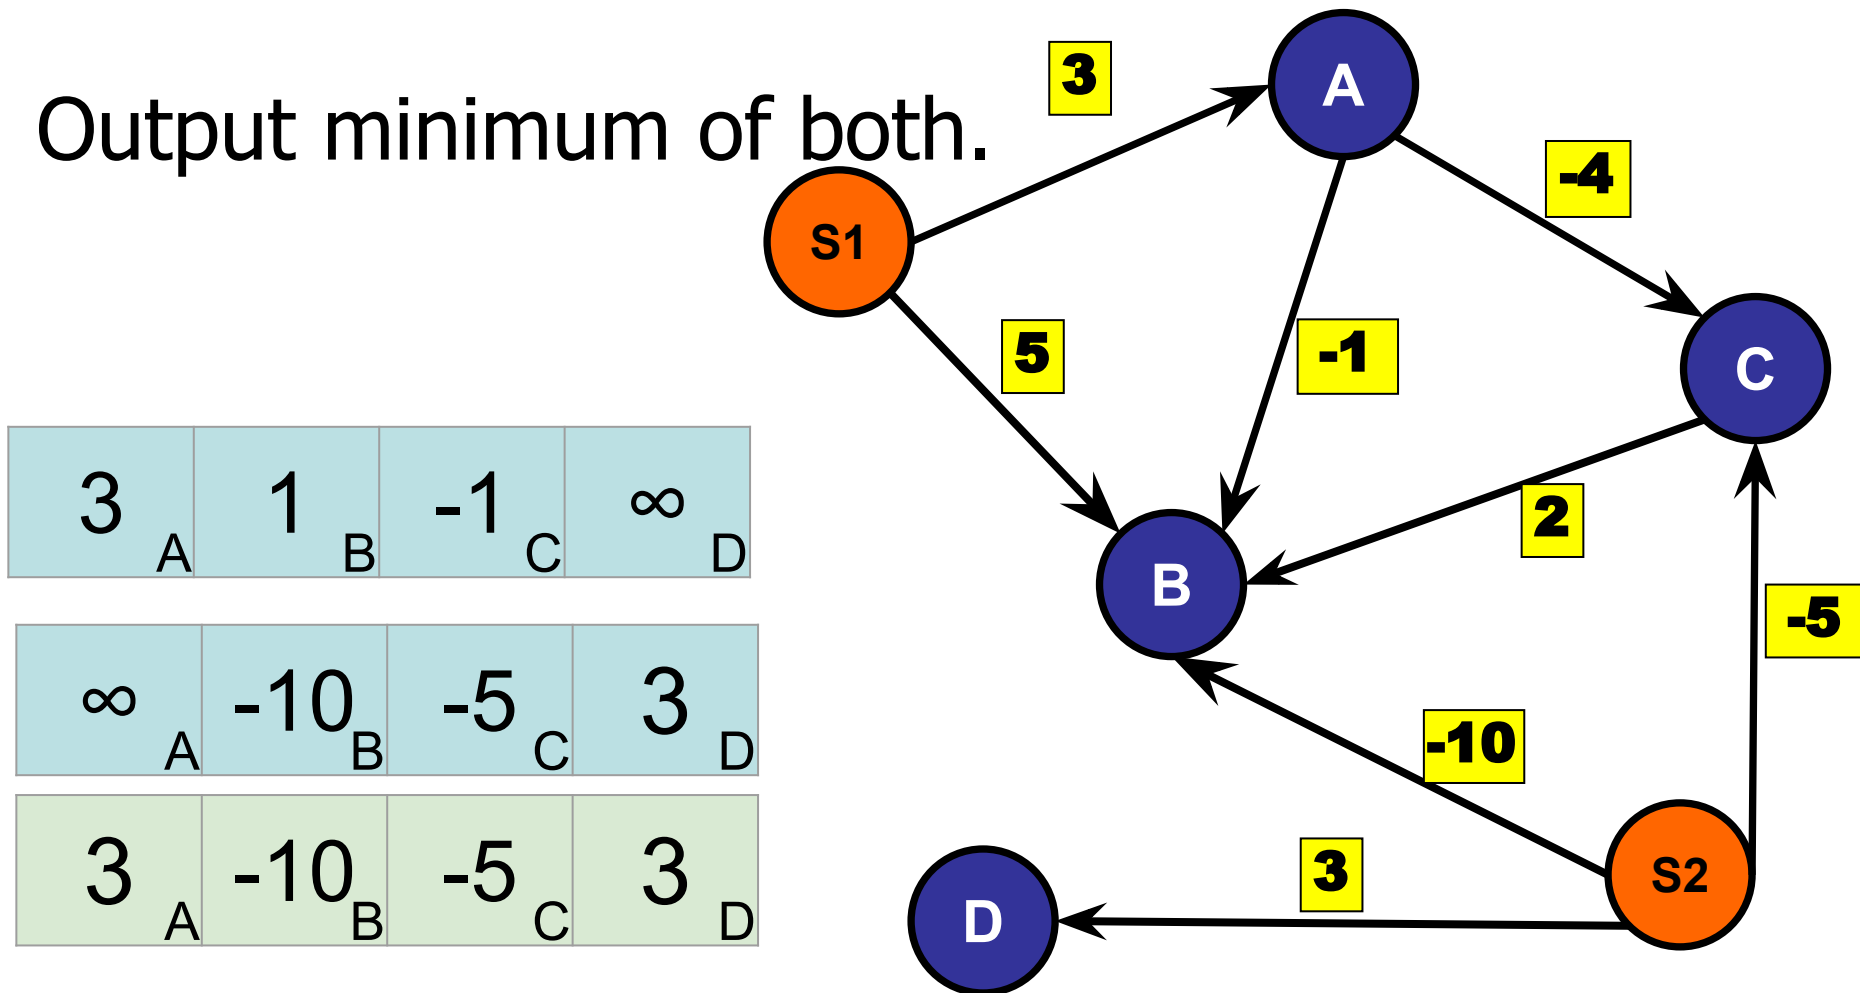
∞	-10	-5	3
A	B	C	D



Technique: Graph Modifications

Obvious solution: Run SSSP once from S1.
Run SSSP again from S2.

Output minimum of both.




Technique: Graph Modifications

But if we had t sources, this means running t copies of SSSP.

3	1	-1	∞
A	B	C	D

∞	-10	-5	3
A	B	C	D

3	-10	-5	3
A	B	C	D



Technique: Graph Modifications


But if we had t sources, this means running t copies of SSSP.

How do we do better?

3	1	-1	∞
A	B	C	D

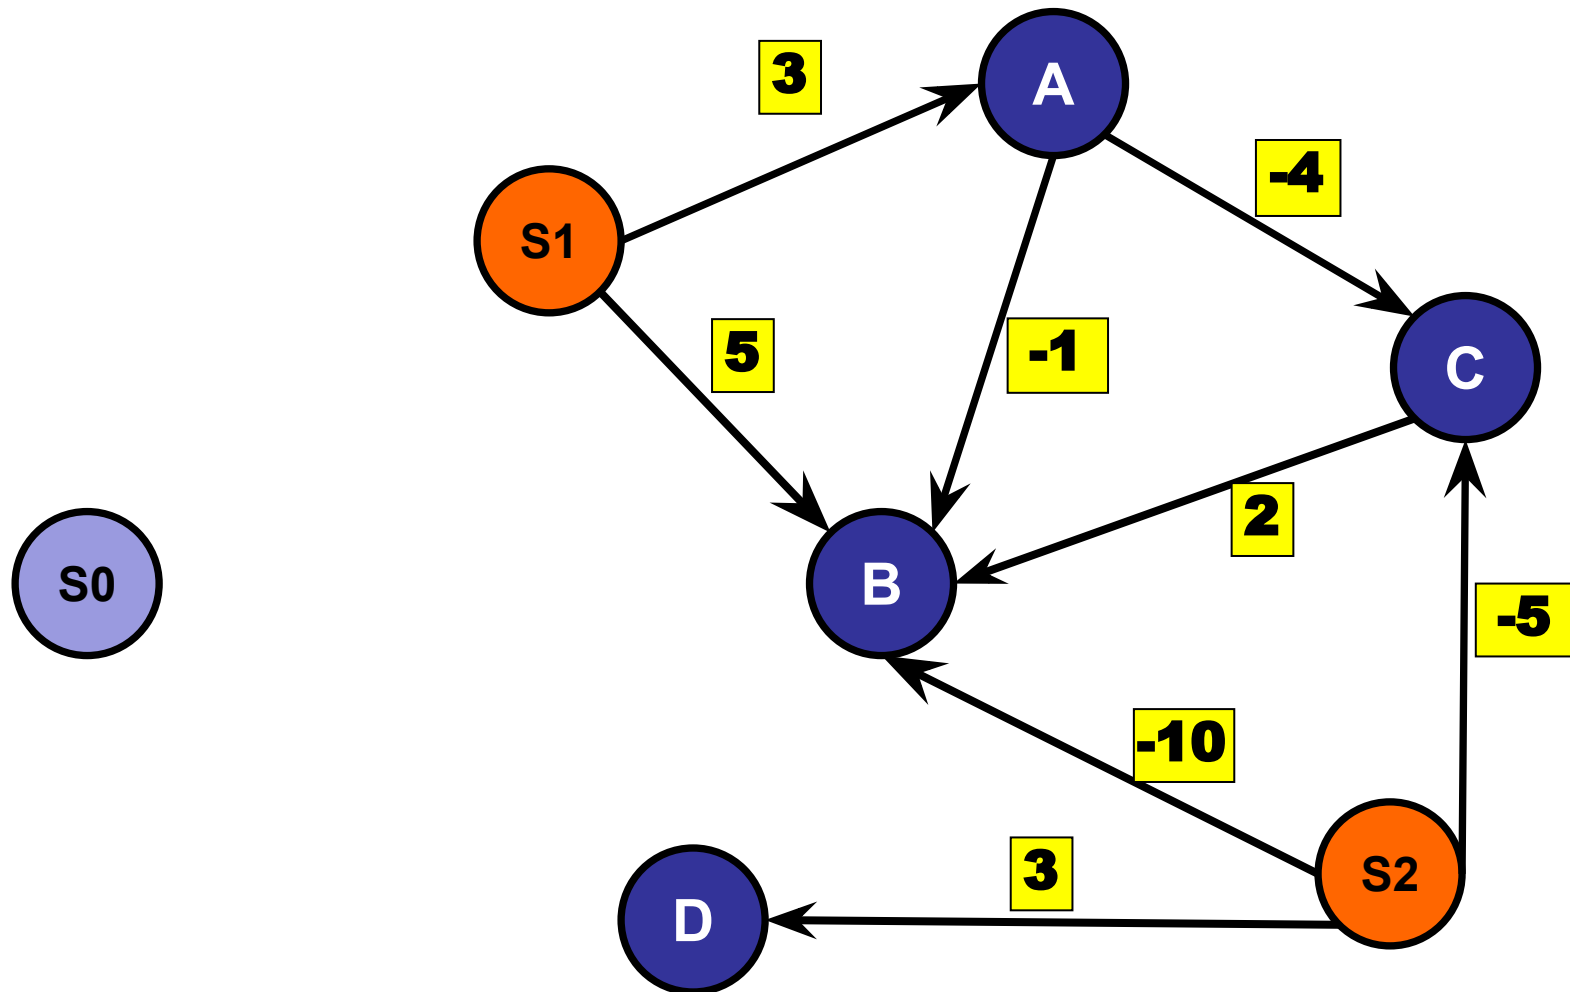
∞	-10	-5	3
A	B	C	D

3	-10	-5	3
A	B	C	D




Technique: Graph Modifications

Idea: Make a false SINGLE source, s0.




Technique: Graph Modifications

Idea: Point s_0 to all sources, with edge costing 0



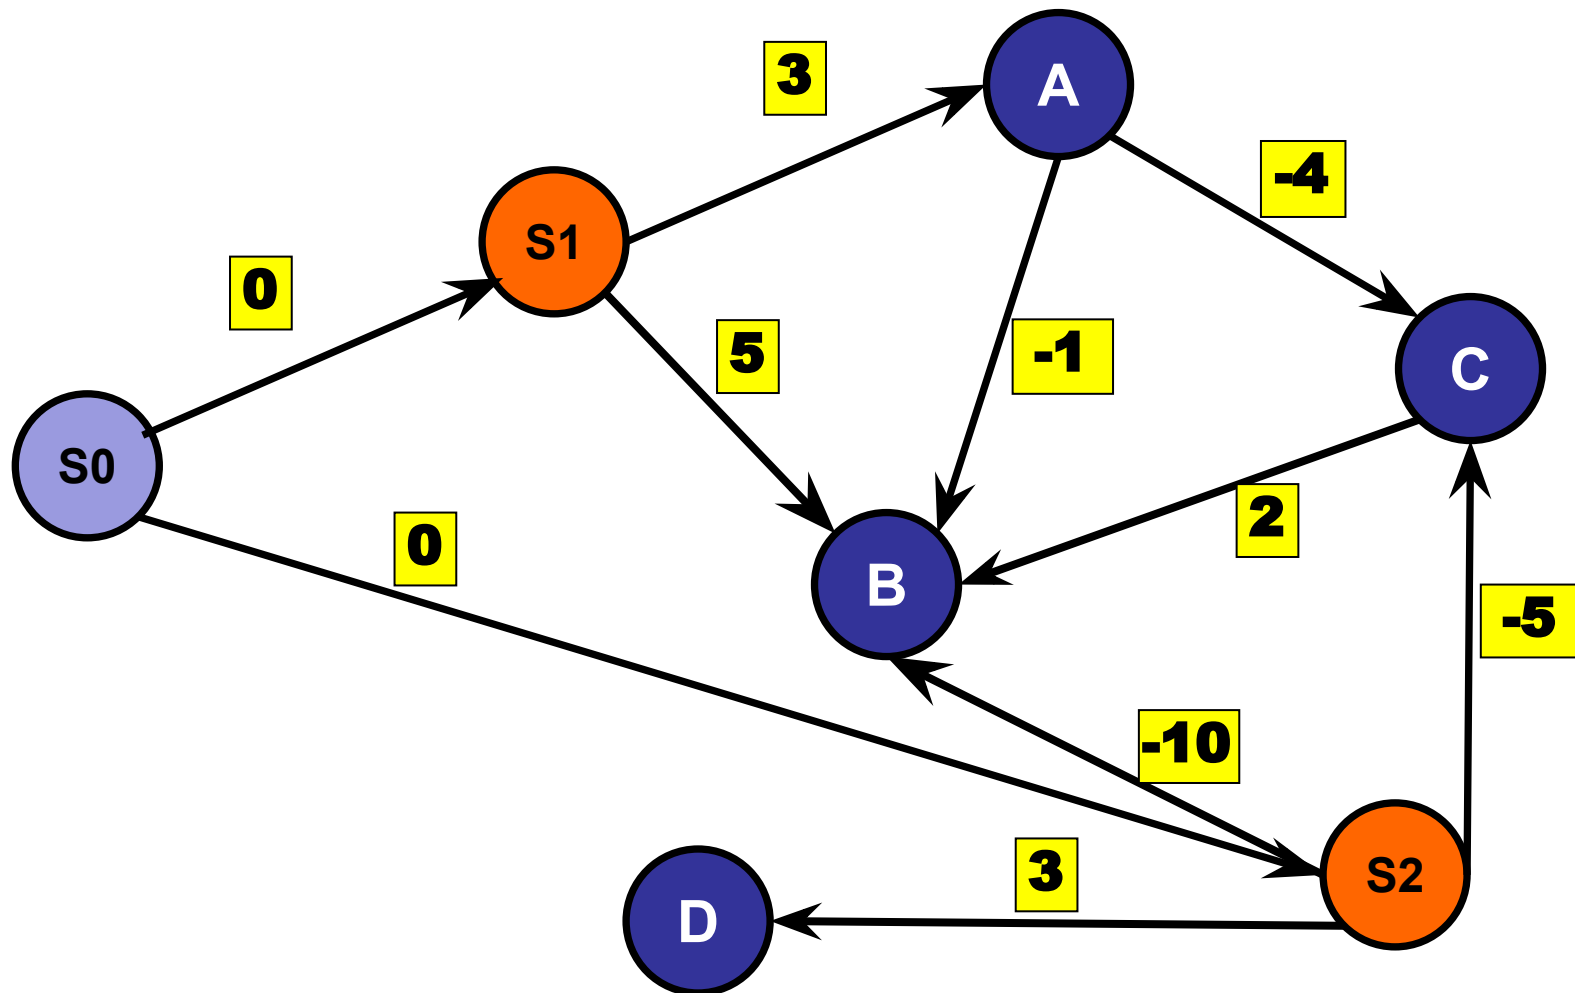
Technique: Graph Modifications

Idea: Run a single copy of SSSP from s_0 .



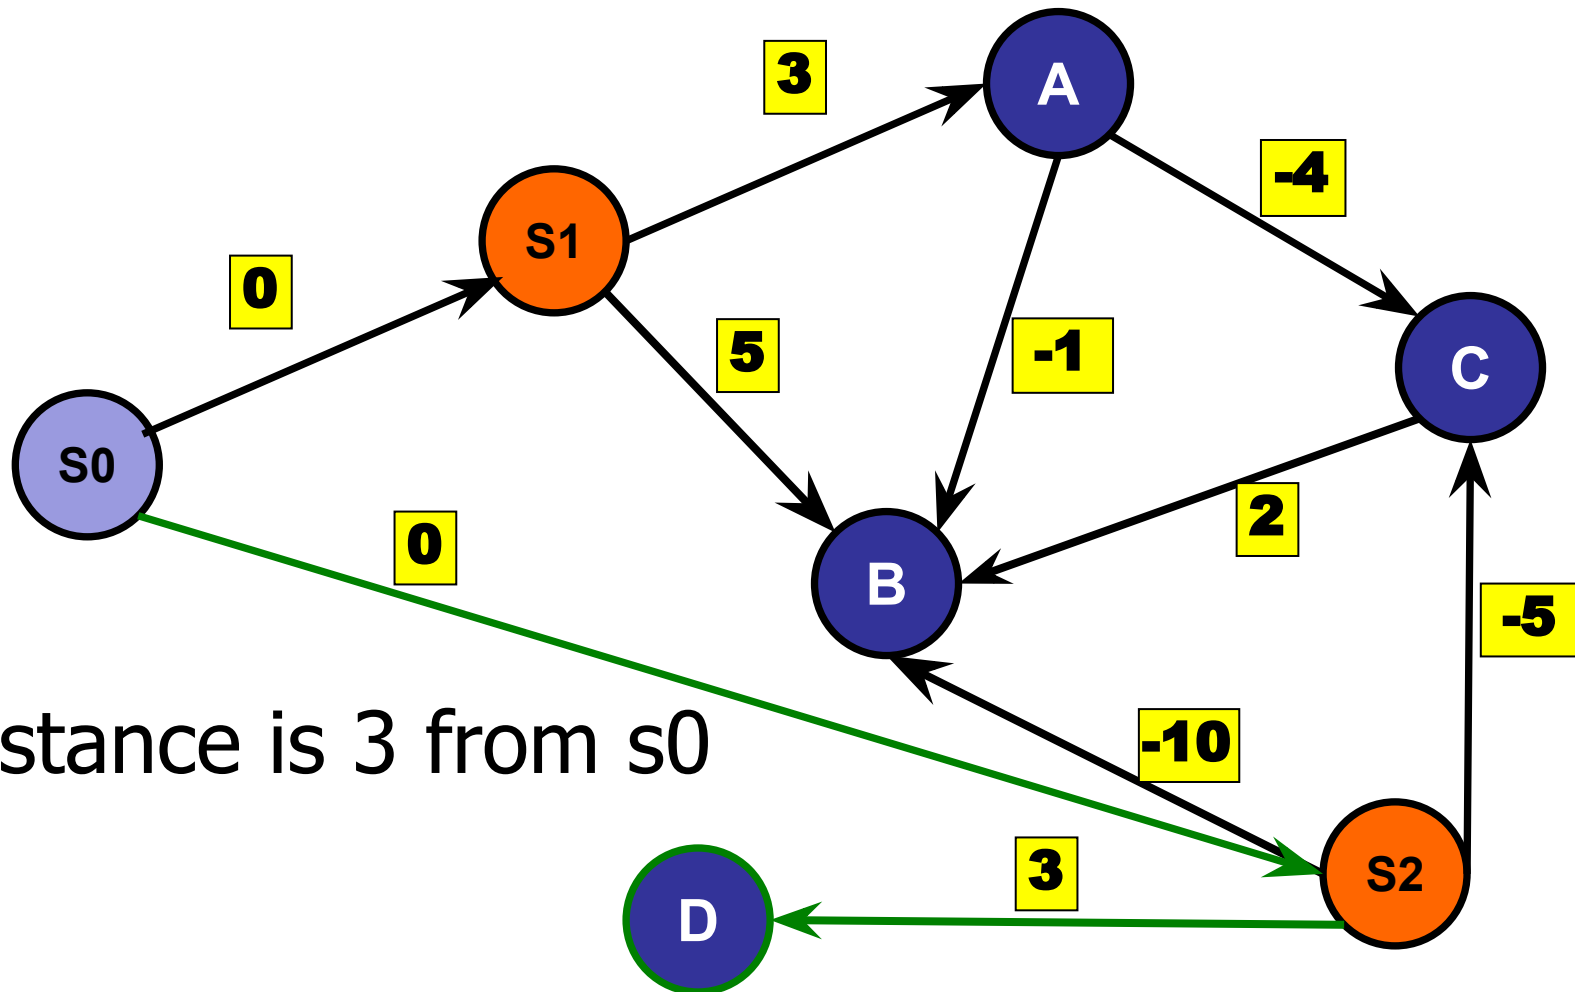
Technique: Graph Modifications

The shortest path from s0 must go through one of the original sources.



Technique: Graph Modifications

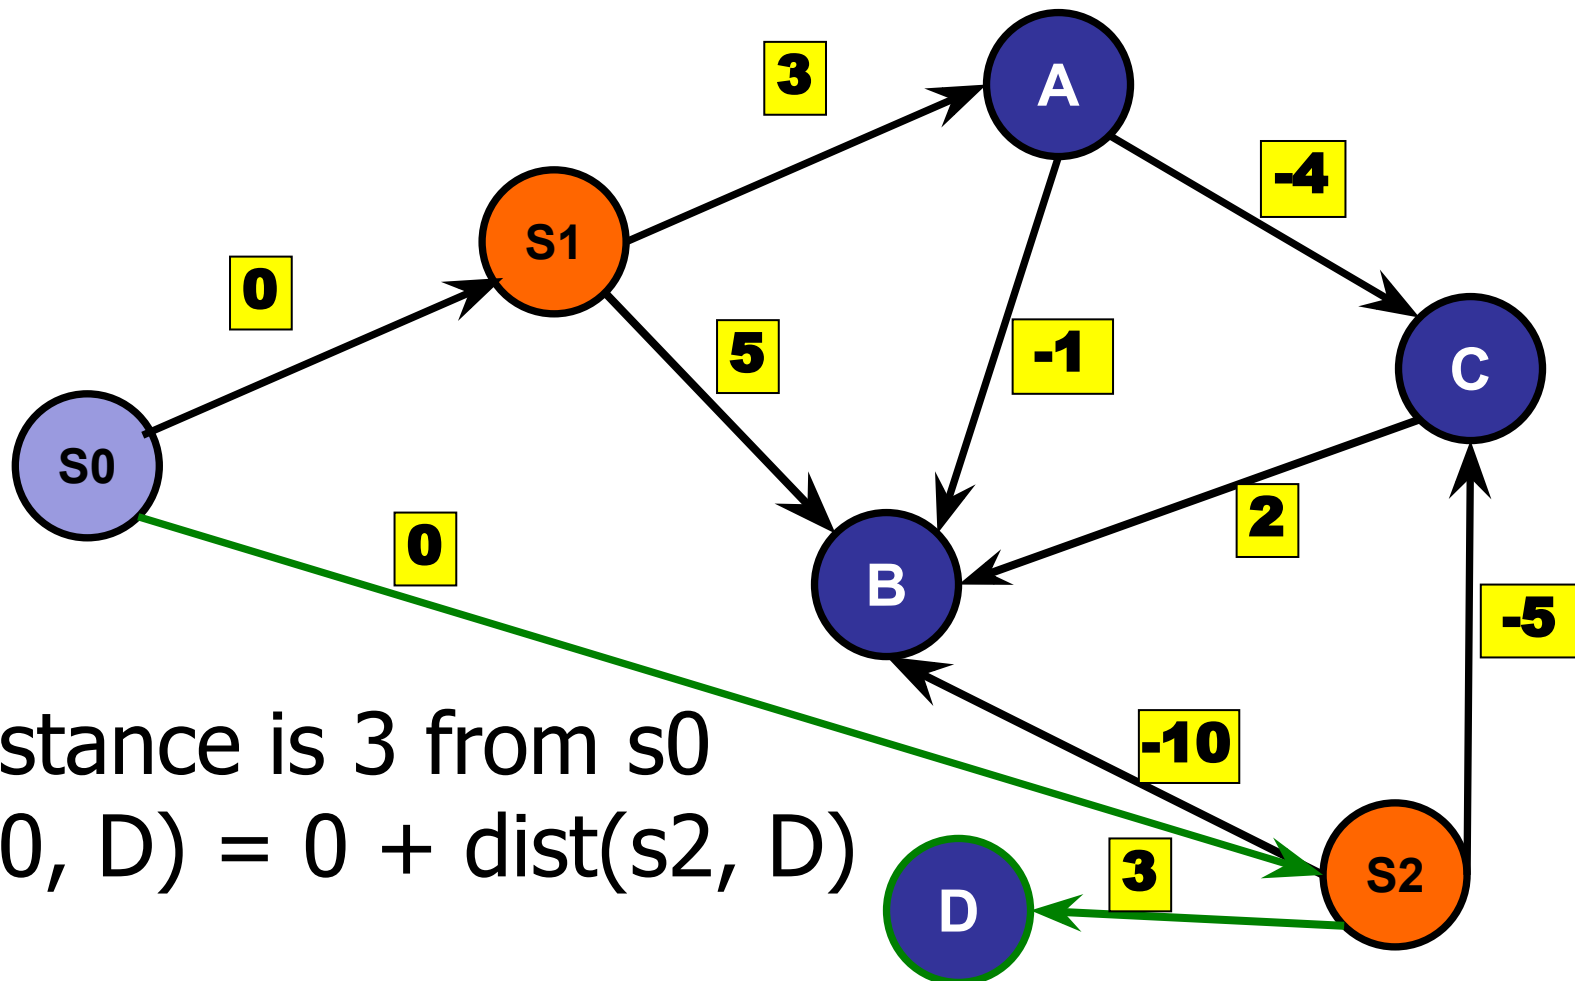
The shortest path from s_0 must go through one of the original sources.



E.g.
D's distance is 3 from s_0

Technique: Graph Modifications

The shortest path from s_0 must go through one of the original sources.



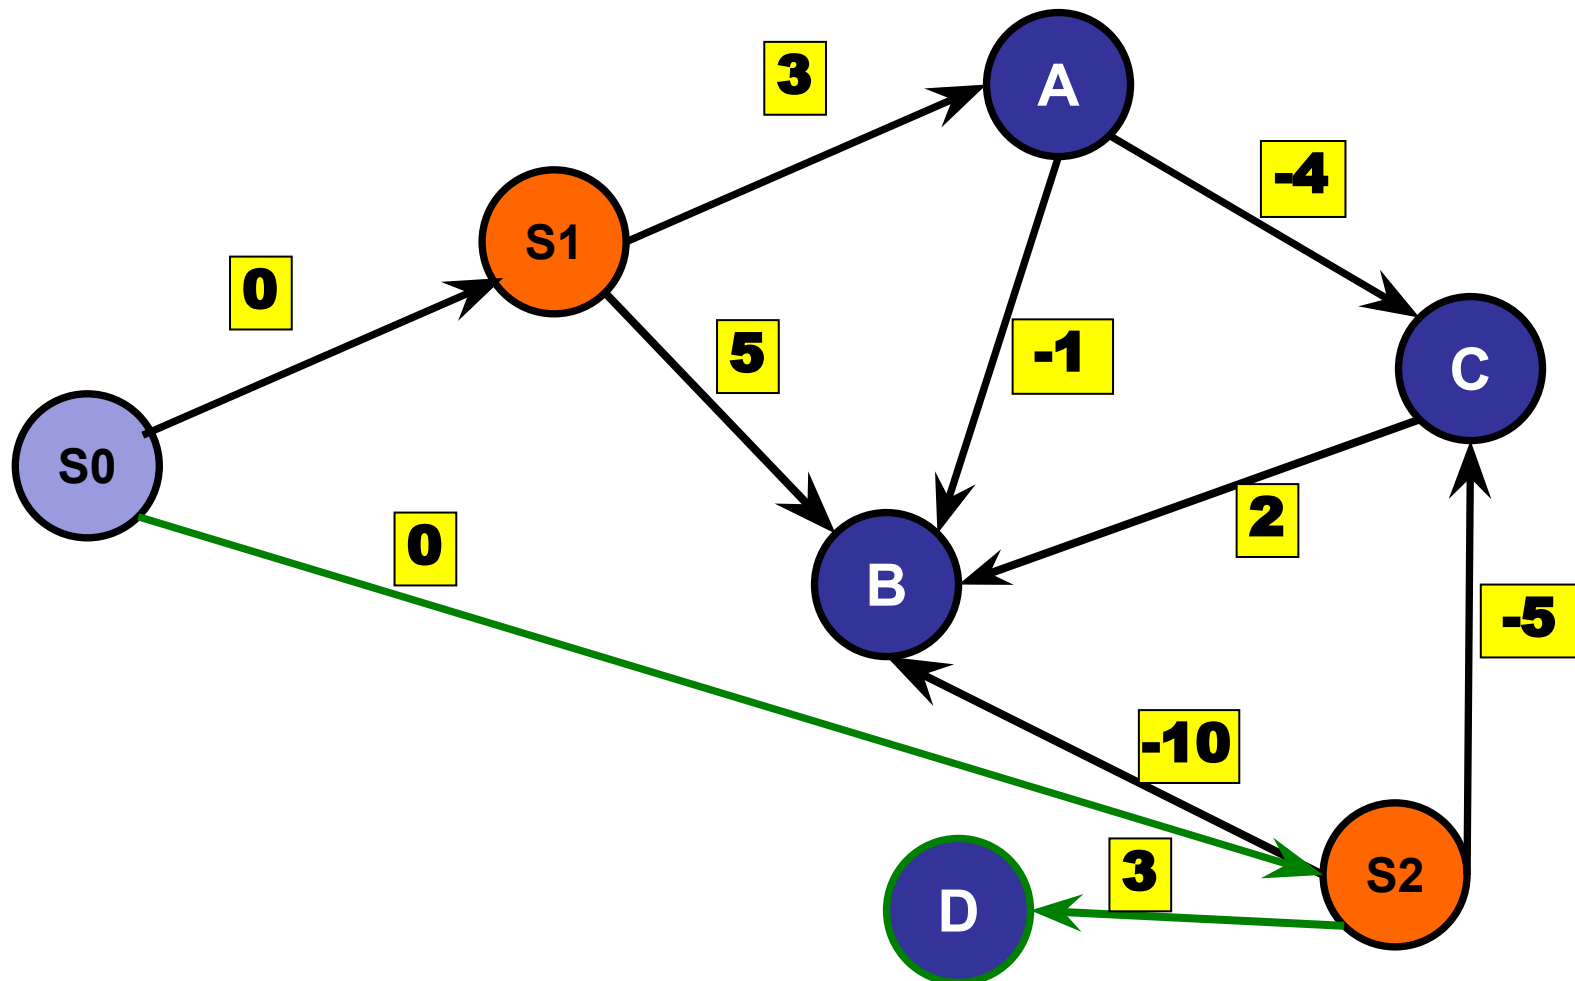
E.g.

D's distance is 3 from s_0

$$\text{dist}(s_0, D) = 0 + \text{dist}(s_2, D)$$

Technique: Graph Modifications

This solves the problem!



Technique: Graph Modifications

What about if we want shortest path that takes at **exactly** k edges?


Technique: Graph Modifications

What about if we want shortest path that takes at **exactly** k edges?

Now we can't just run SSSP because we don't know how many edges the shortest path takes.


Technique: Graph Modifications

E.g. Shortest from A to C using exactly ? edge.




Technique: Graph Modifications

E.g. Shortest from A to C using exactly 1 edge.
= impossible




Technique: Graph Modifications

E.g. Shortest from A to C using exactly 2 edges.
= 4



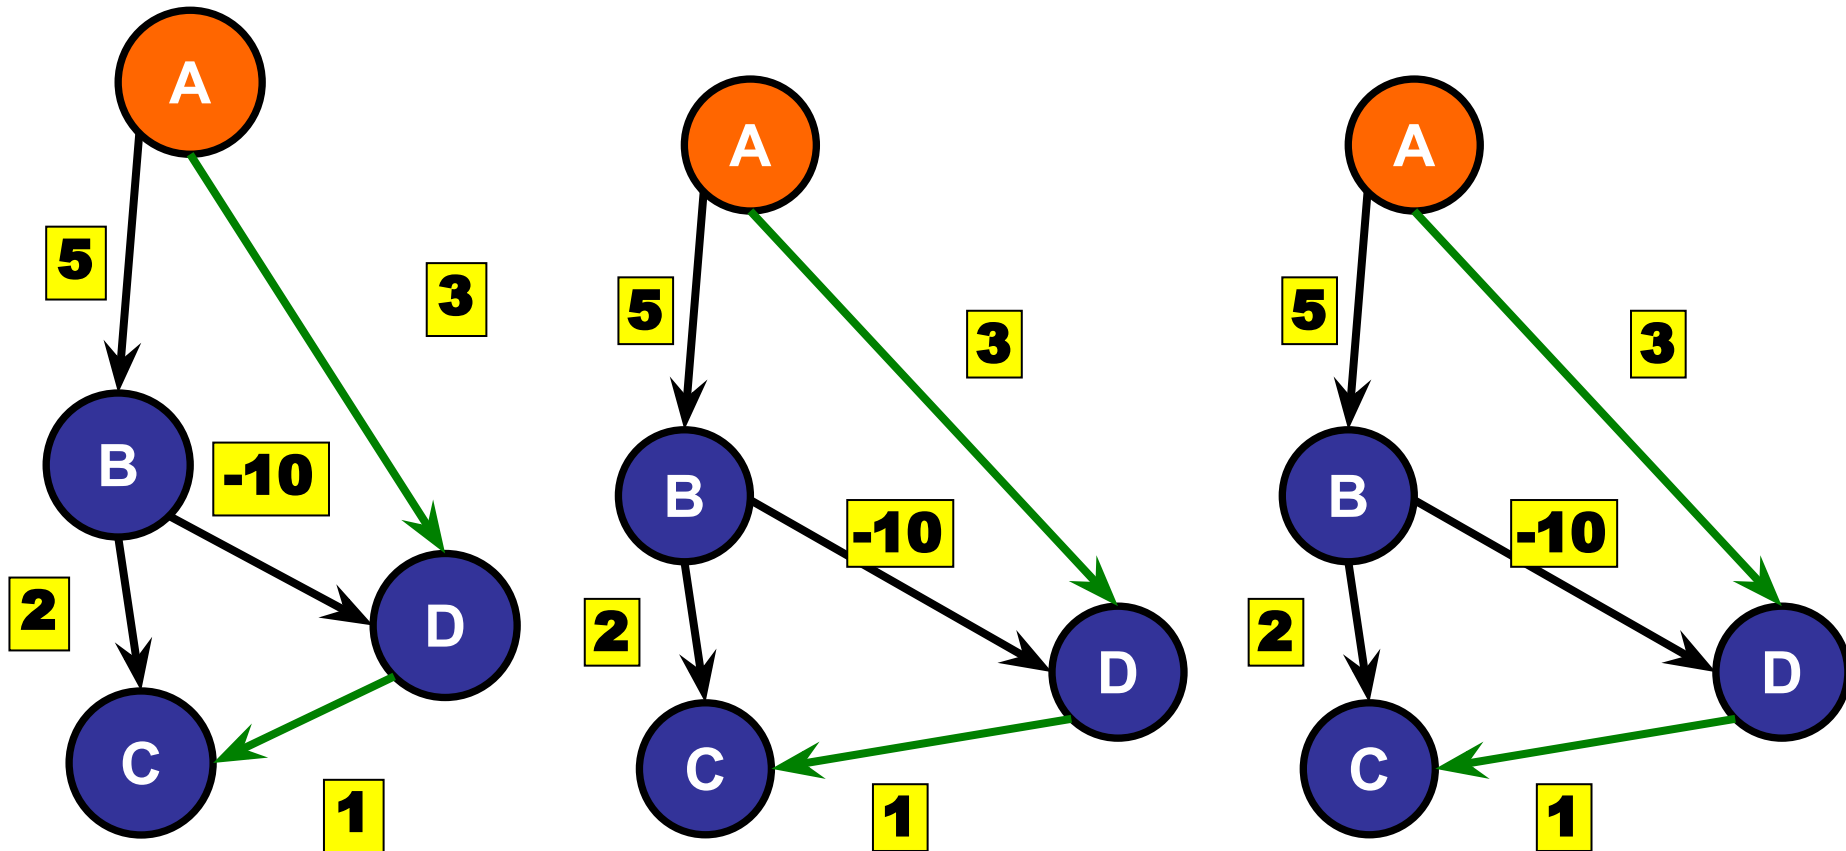
Technique: Graph Modifications

E.g. Shortest from A to C using exactly 3 edges.
 $= 5 - 10 + 1 = -4$



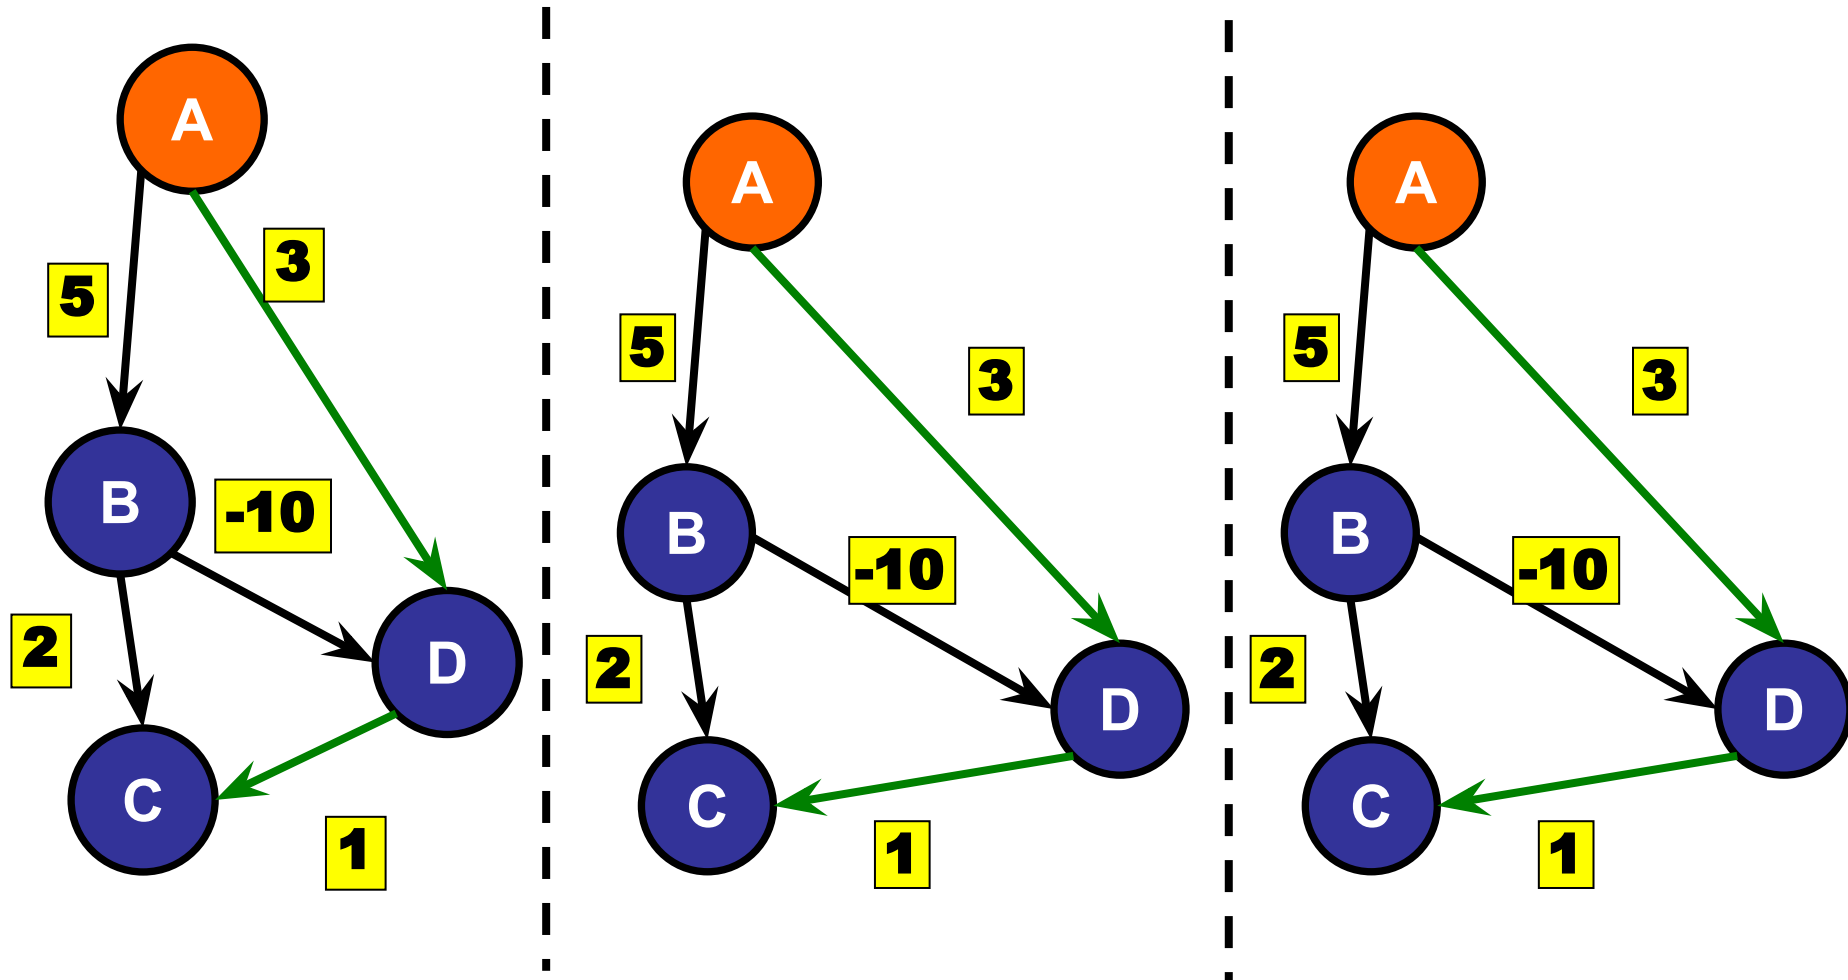
Technique: Graph Modifications

Idea, what happens if we copied the graph $k + 1$ times? E.g. $k = 2$




Technique: Graph Modifications

Call each copy a layer:




Technique: Graph Modifications

Now if originally, A goes to D,
then A in layer 0 goes to D in layer 1.



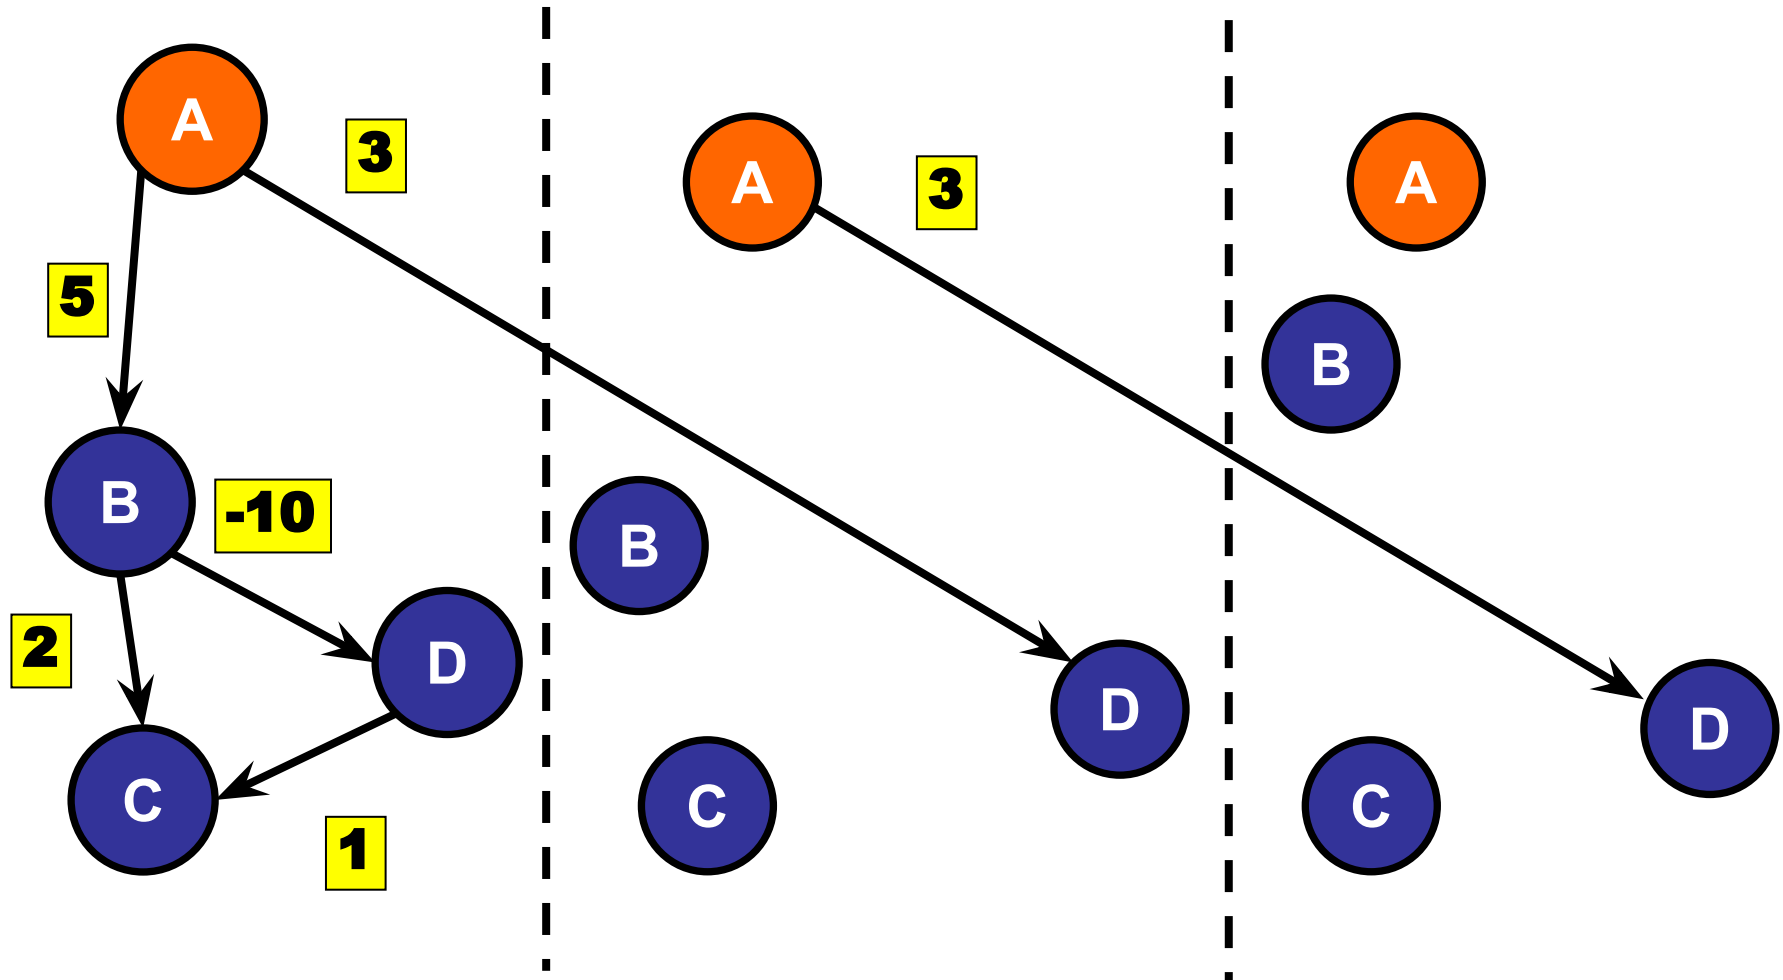
Technique: Graph Modifications

Now if originally, A goes to D,
then A in layer 0 goes to D in layer 1.




Technique: Graph Modifications

Now if originally, A goes to D, similarly then A in layer 1 goes to D in layer 2.



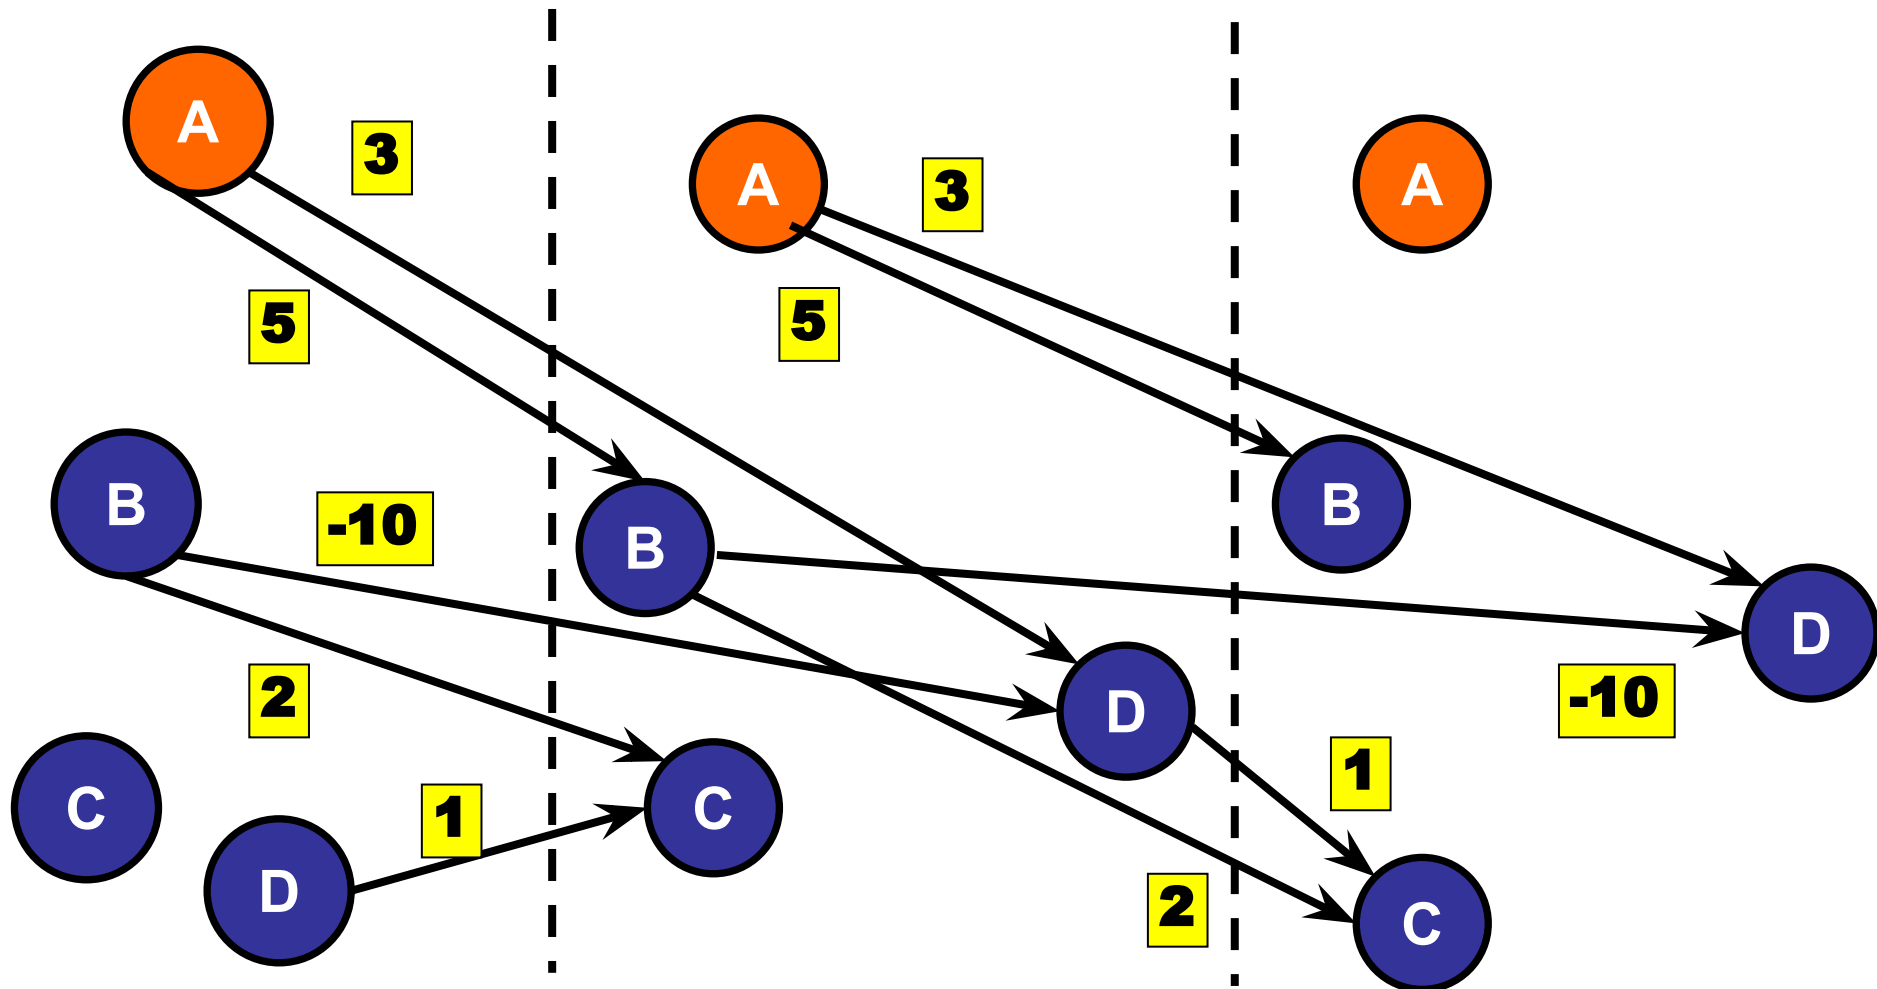
Technique: Graph Modifications

Do this for all edges: if (u, v) is an edge,
Then draw edge from u from layer i to v from layer $i + 1$.



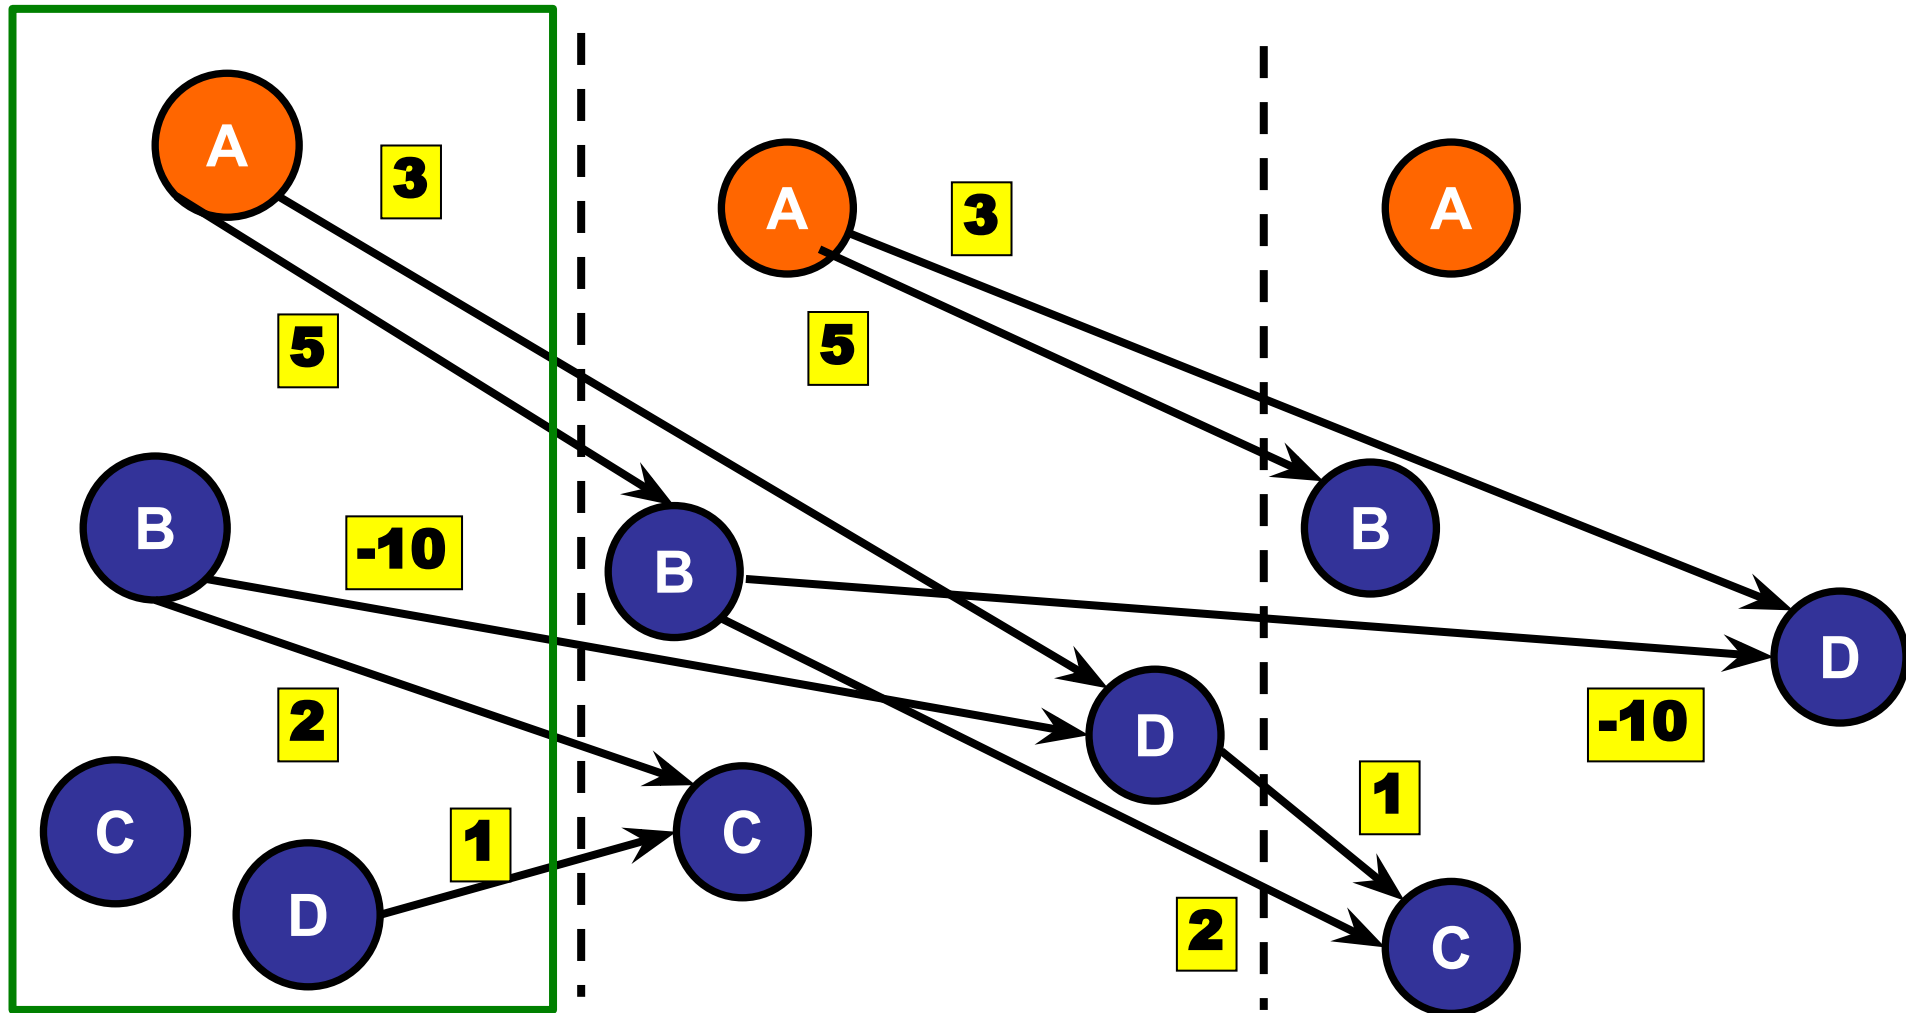
Technique: Graph Modifications

Intuition: If we are on the i th layer, we have taken exactly i steps in the original graph.




Technique: Graph Modifications

Intuition: If we are on the i th layer, we have taken exactly i steps in the original graph. 0th layer = 0 steps



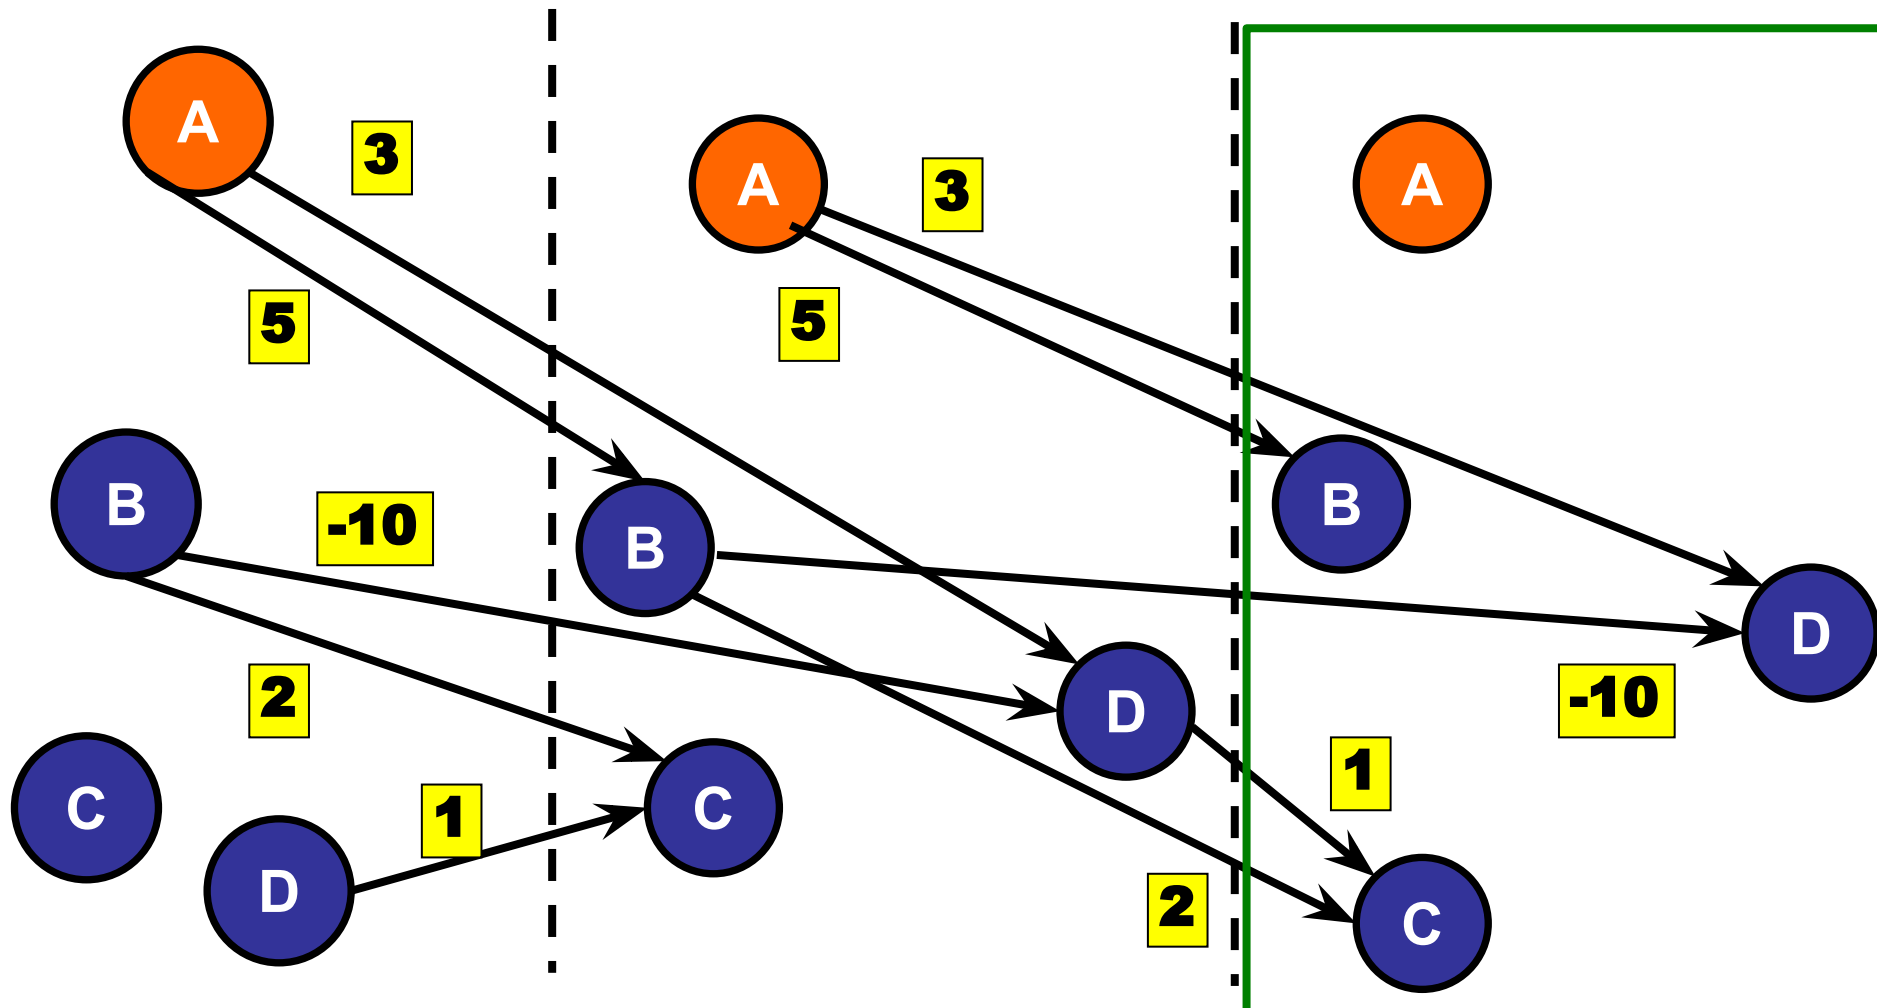
Technique: Graph Modifications

Intuition: If we are on the i th layer, we have taken exactly i steps in the original graph. 1st layer = 1 step



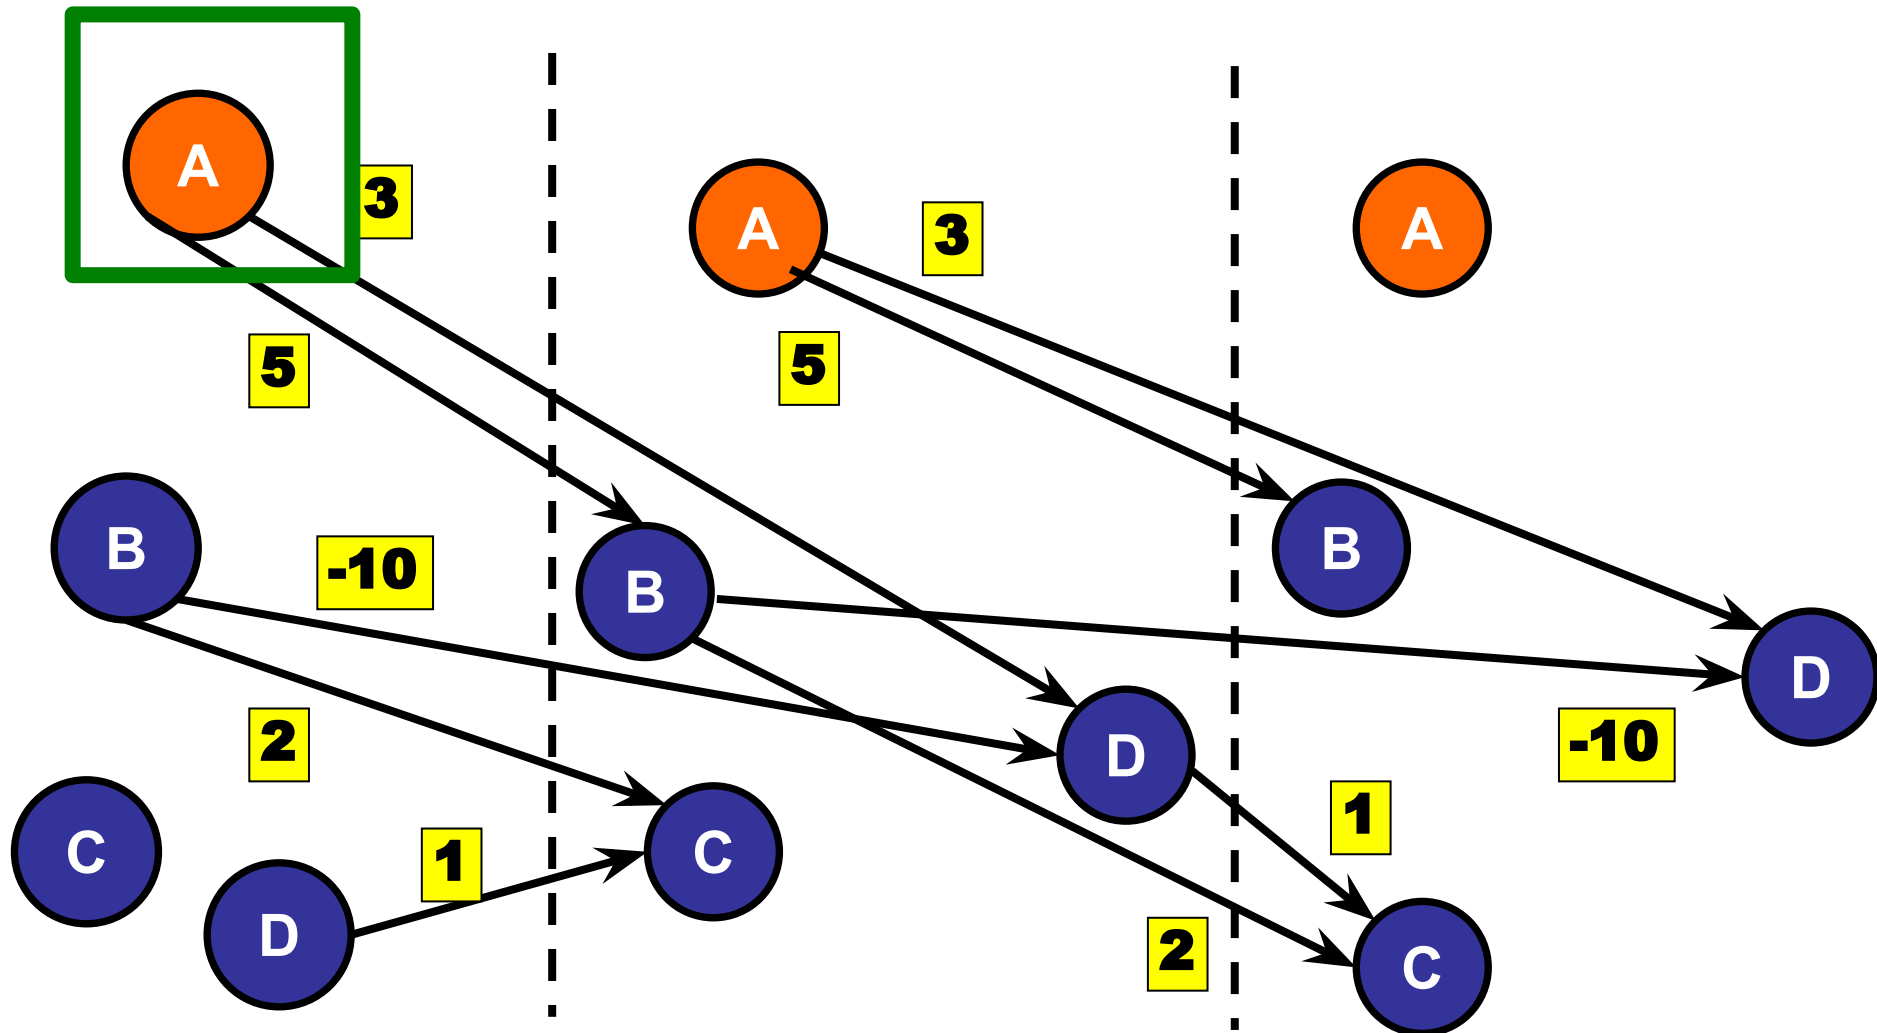
Technique: Graph Modifications

Intuition: If we are on the i th layer, we have taken exactly i steps in the original graph. 2nd layer = 2 steps




Technique: Graph Modifications

If we SSSP from source node in layer 0



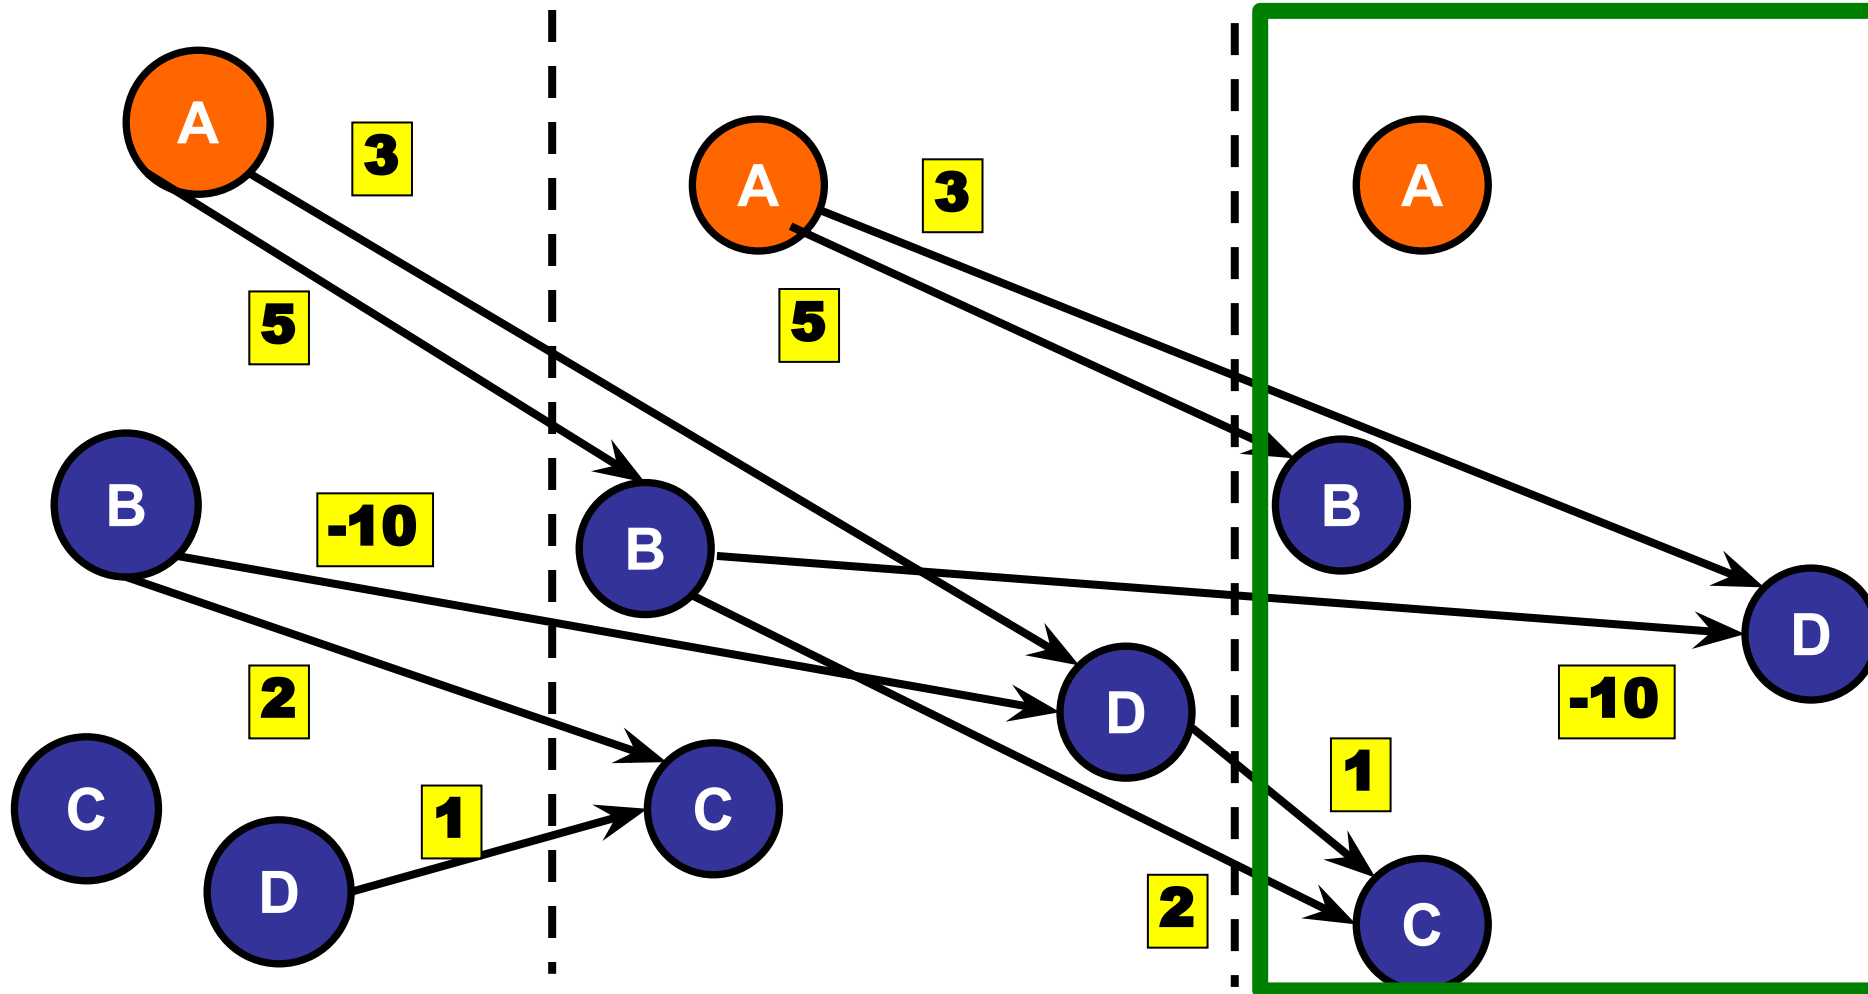
Technique: Graph Modifications

The distances to the nodes in the last layer all took EXACTLY 2 steps.



Technique: Graph Modifications

This technique of creating layers of the same graph is quite general.



Technique: Graph Modifications

This technique of creating layers of the same graph is quite general.

Effectively, we have created a *state space*. Where each node represents a kind of state.

Technique: Graph Modifications

This technique of creating layers of the same graph is quite general.

Effectively, we have created a *state space*. Where each node represents a kind of state.

E.g. Node B in layer 2 represents the state of “taking exactly 2 steps from source node to B.”

Technique: Graph Modifications

This technique of creating layers of the same graph is quite general.

Effectively, we have created a *state space*. Where each node represents a kind of state.

E.g. Node B in layer 2 represents the state of “taking exactly 2 steps from source node to B.”

Can encode many things about your traversal through the graph

Technique: Graph Modifications

This technique of creating layers of the same graph is quite general.

Effectively, we have created a *state space*. Where each node represents a kind of state.

E.g. Node B in layer 2 represents the state of “taking exactly 2 steps from source node to B.”

Can encode many things about your traversal through the graph

Bonus Slides: 3AM Things:

So Far:

- **Unweighted** graph
 - BFS
 - $O(V + E)$
- **Weighted** graphs with non-negotiable

Google slides putting words in my mouth.

Today

Single Source Shortest Paths (SSSP):

- Bellman Ford
 - SSSP on negative edge graphs
 - Negative Cycle Detection

Some graph techniques

Next Week:

MSTs and Union Find